

# TensorFlow 2.0 Tutorial: Part #2

High-Level APIs (Sequential, Functional, and Model Subclassing) and more!



Iran University of Science and Technology (IUST)  
Department of Computer Engineering



Notebook URL:

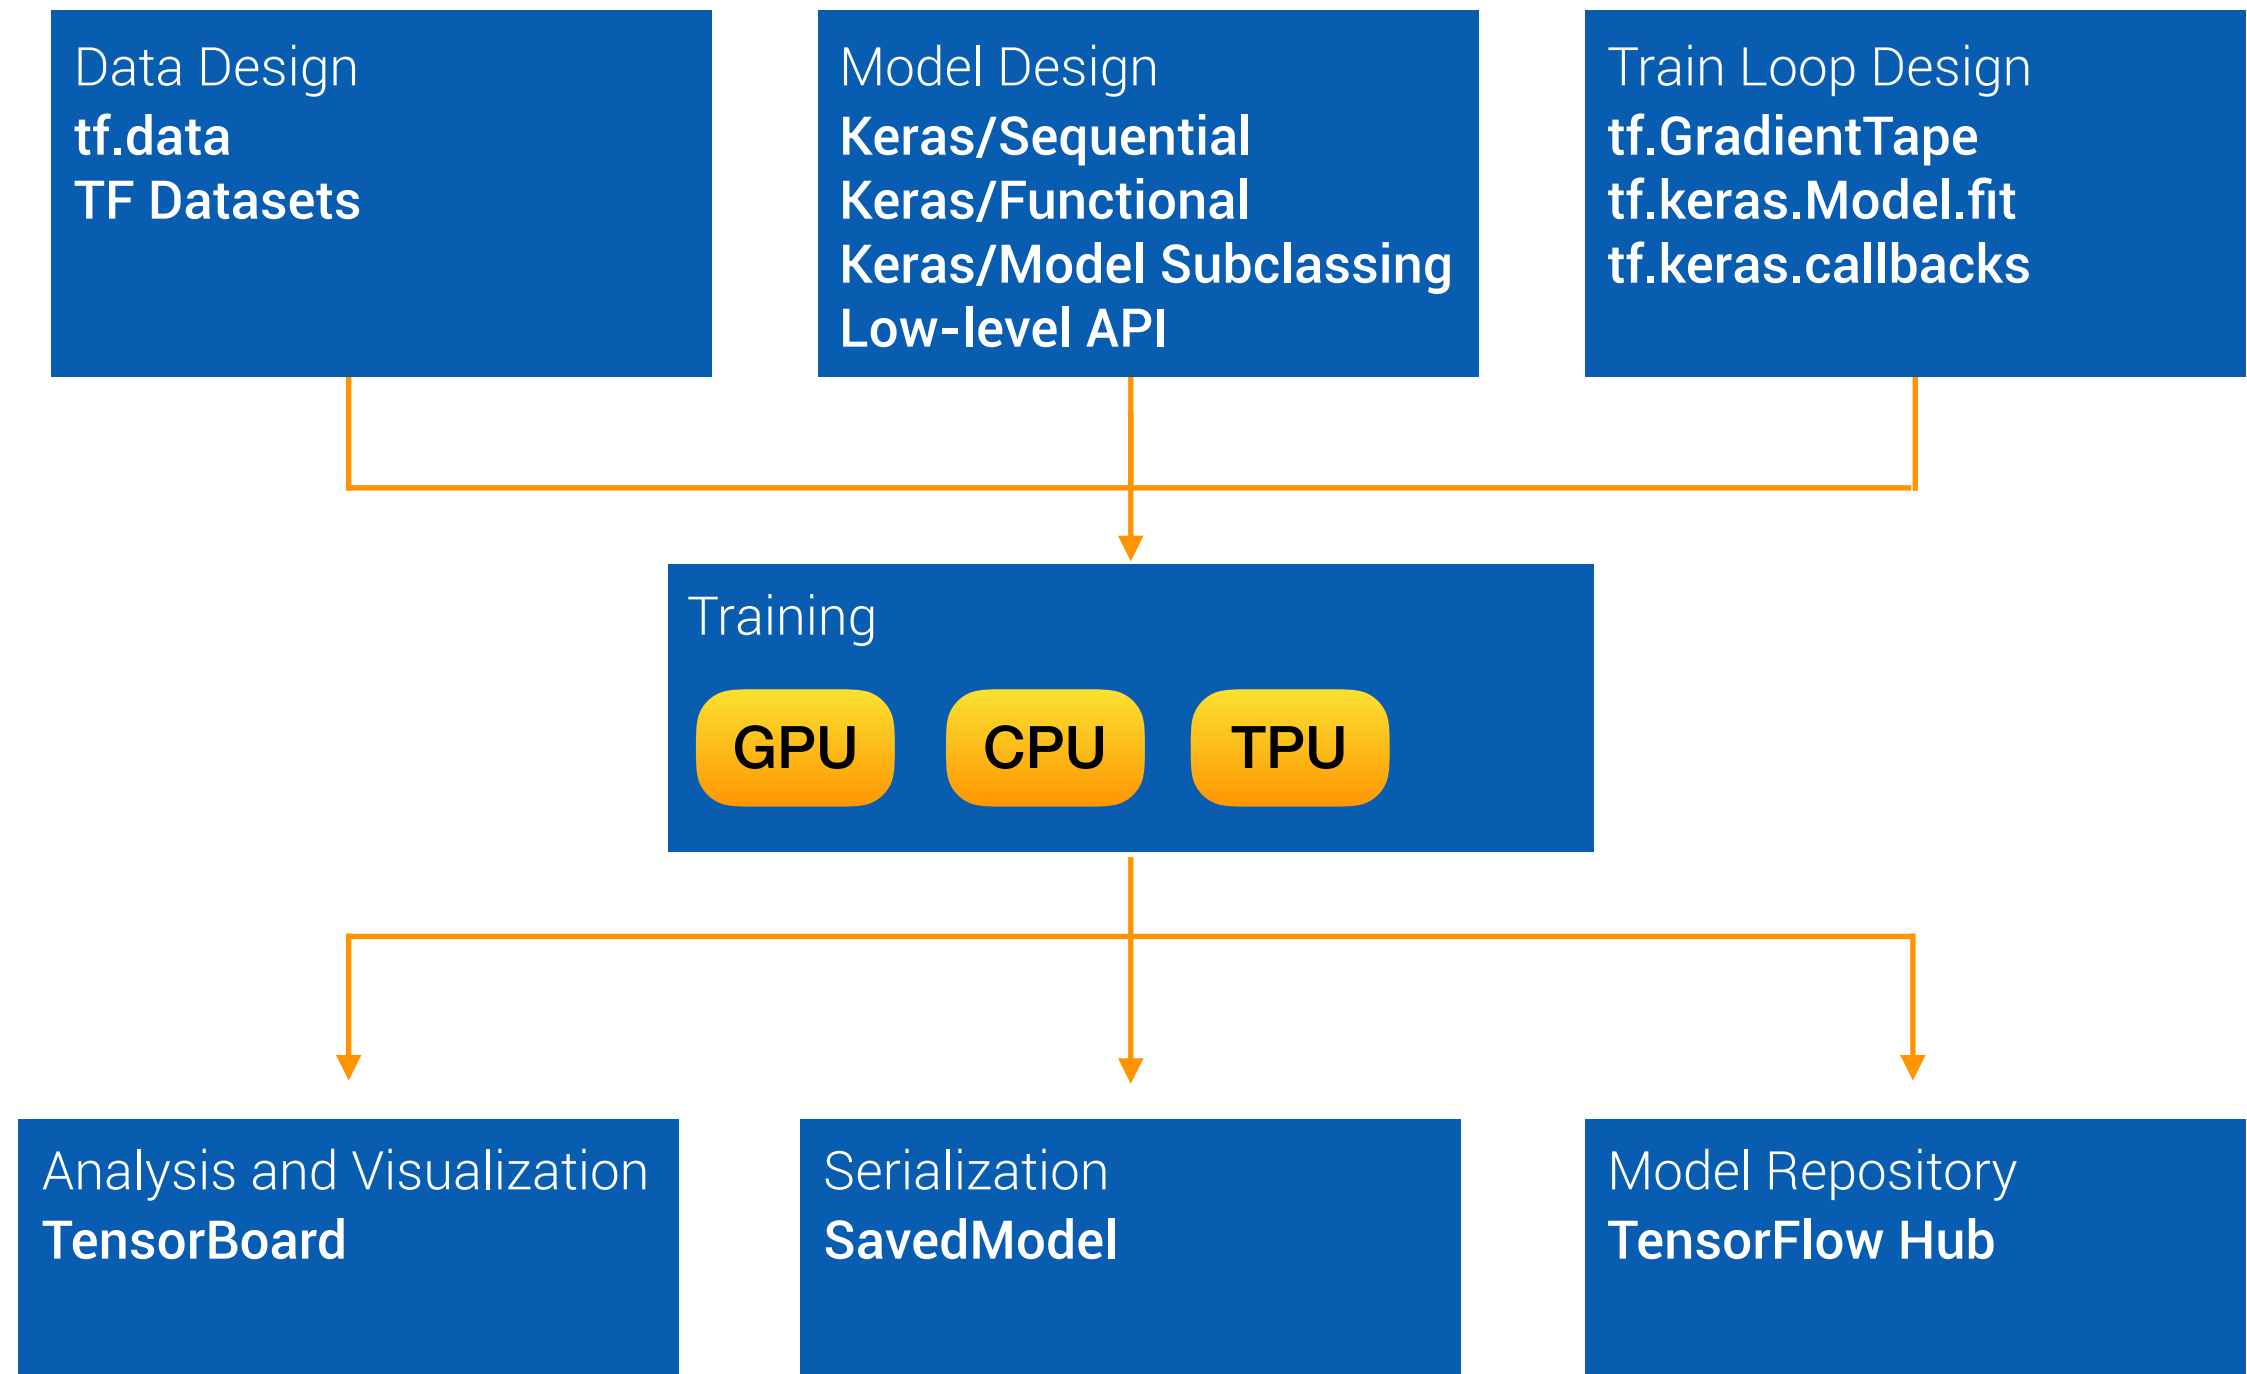
[https://github.com/kazemnejad/tensorflow-2-tutorial/blob/master/part\\_02.ipynb](https://github.com/kazemnejad/tensorflow-2-tutorial/blob/master/part_02.ipynb)

Slides URL:

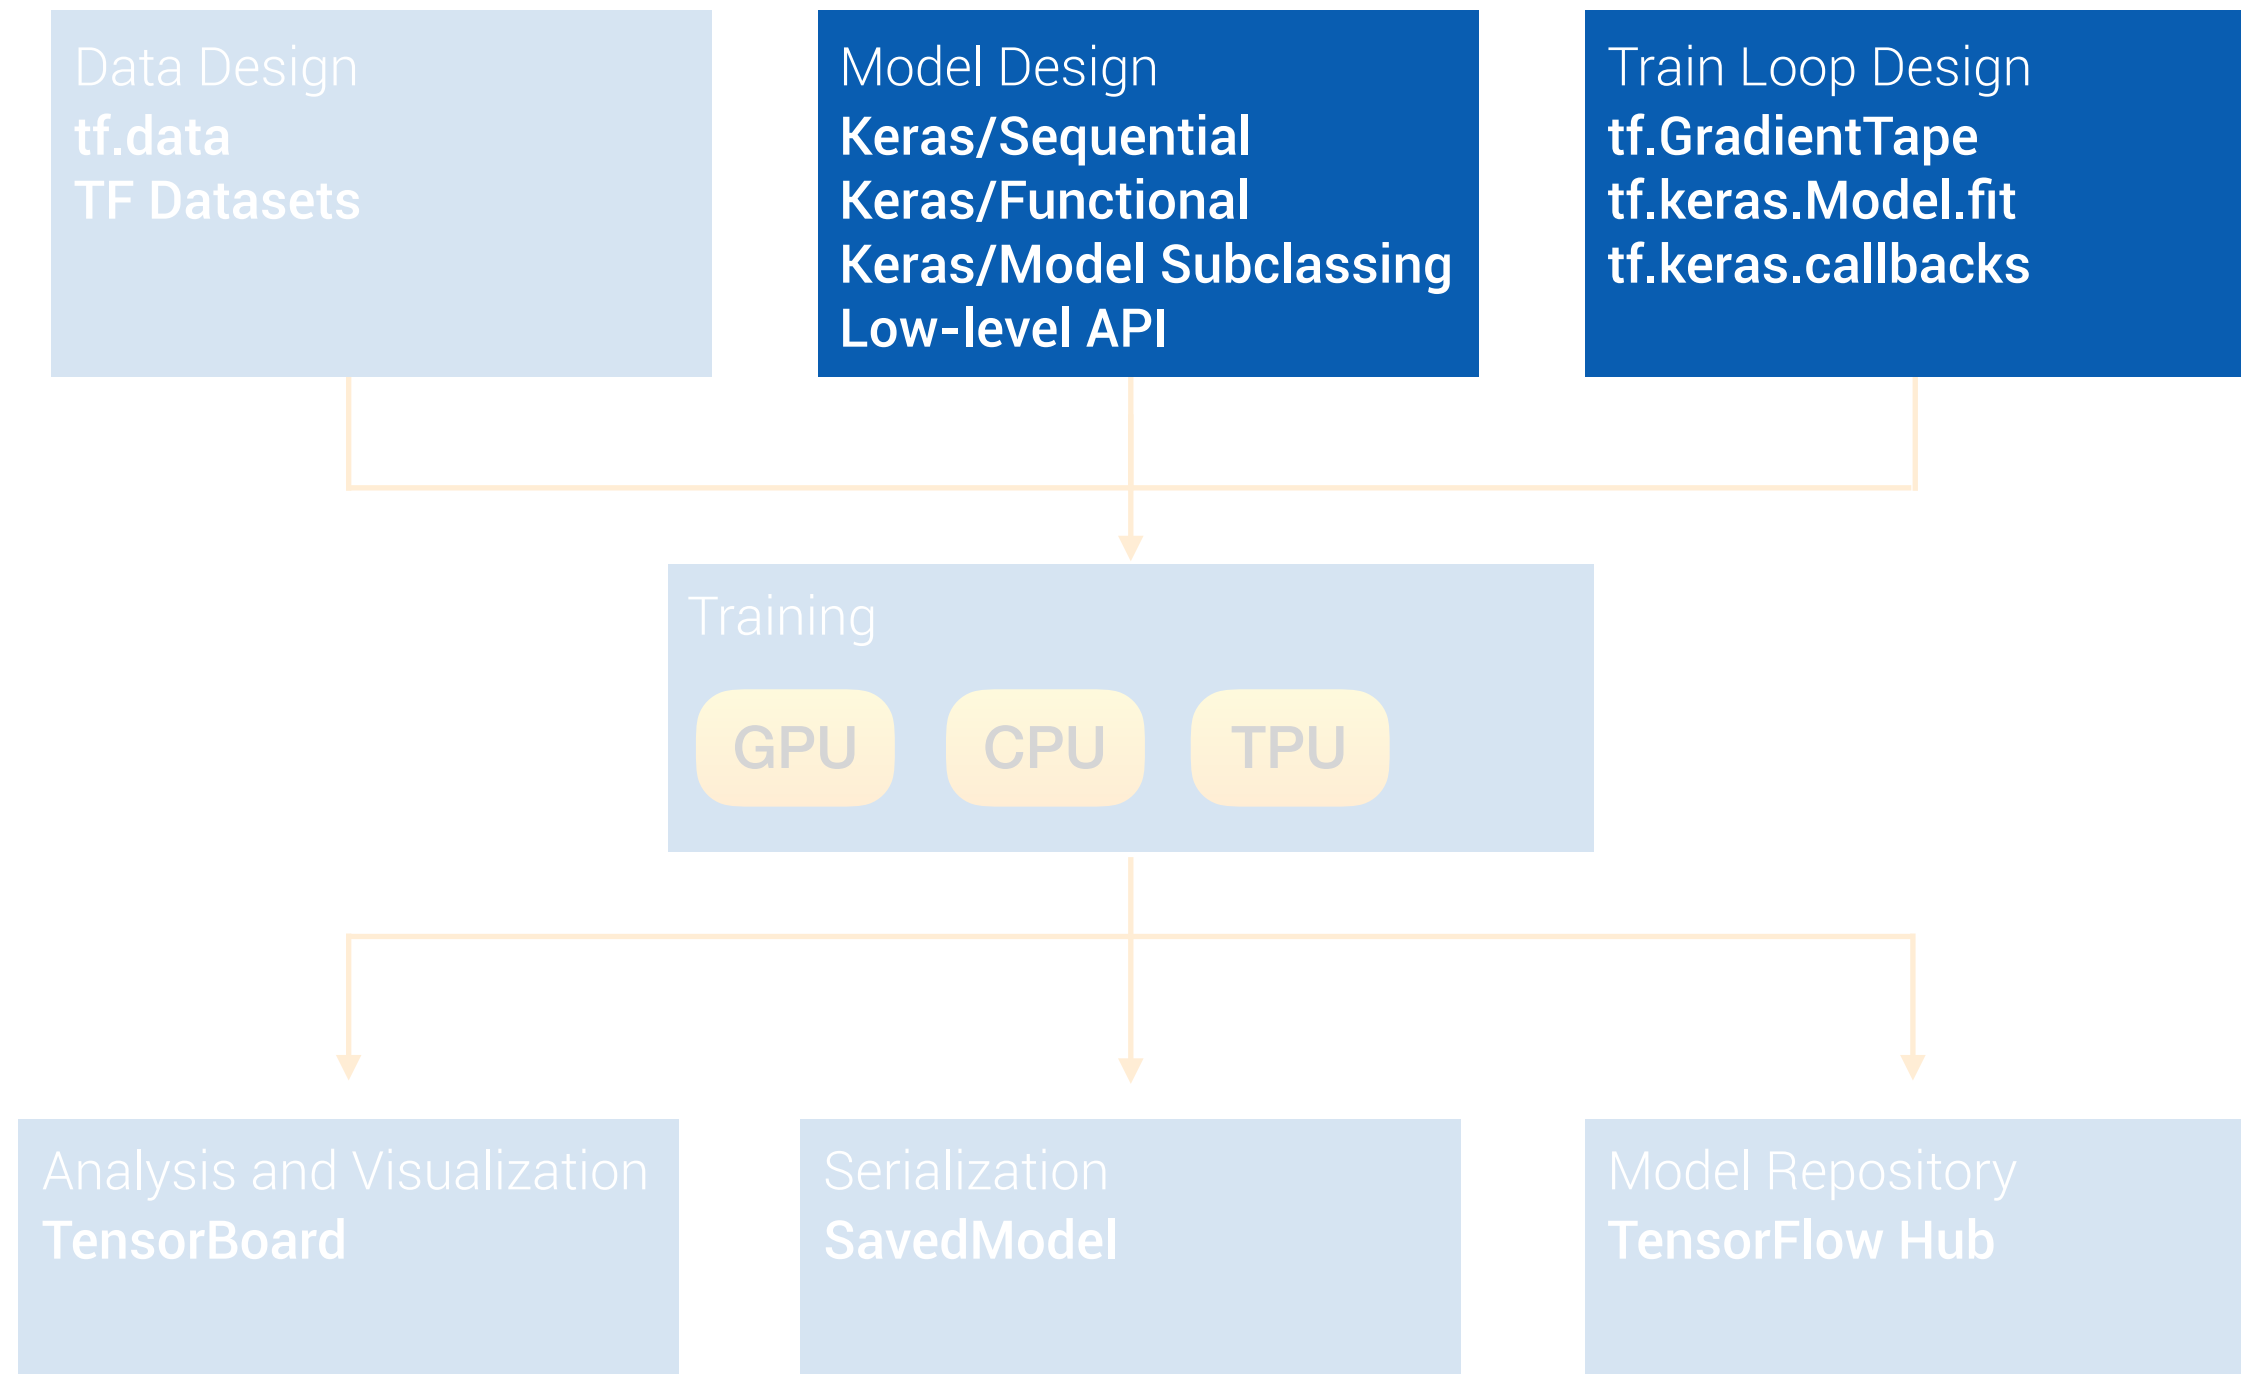
[https://github.com/kazemnejad/tensorflow-2-tutorial/blob/master/part\\_02\\_slides.pdf](https://github.com/kazemnejad/tensorflow-2-tutorial/blob/master/part_02_slides.pdf)



# TensorFlow Overview



# TensorFlow Overview



## Package

**keras.\***

**vs**

**tf.keras.\***

- **tf.keras** is a re-implementation of the Keras API.
- **tf.keras** has better Integration with rest of the framework.
- Distributed training is much easier in **tf.keras**.
- **tf.keras** supports Eager execution (dynamic graph).
- There is no one-to-one relation. However, most of the useful stuffs are also present in TensorFlow.

# Model Design

- Keras **Sequential** API
  - + standard layers
  - + custom layers, losses, and metrics
  - Stack of layers
  - For Simple models



# TensorFlow Higher Level APIs

- Keras API (`tf.keras.*`)
  - Engine
    - **Base Layer, Base Model, Sequential**
  - Layers (various subclasses of Base Layer)
  - Losses, Metrics
  - Callbacks
  - Optimizers
  - Regularizes, Constraints

# TensorFlow Higher Level APIs

- `tf.Module()` (Base neural network module class)
- Keras API (`tf.keras.*`)
  - Engine
    - **Base Layer, Base Model, Sequential**
  - Layers (various subclasses of Base Layer)
  - Losses, Metrics
  - Callbacks
  - Optimizers
  - Regularizes, Constraints



```
import tensorflow as tf
from tensorflow.keras import layers

model = tf.keras.Sequential()
model.add(layers.Dense(32, activation='relu', input_shape=(784,)))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

```
import tensorflow as tf
from tensorflow.keras import layers

model = tf.keras.Sequential([
    layers.Dense(32, activation='relu', input_shape=(784,)),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
```





What does a Layer do?

# What does a Layer do?

- Computation from **a batch** of inputs to **a batch** of outputs.

## What does a Layer do?

- Computation from **a batch** of inputs to **a batch** of outputs.
- Manages **state** (trainable **weights**, non-trainable weights).

## What does a Layer do?

- Computation from **a batch** of inputs to **a batch** of outputs.
- Manages **state** (trainable **weights**, non-trainable weights).
- Tracks **losses** and **metrics**.



## What does a Layer do?

- Computation from **a batch** of inputs to **a batch** of outputs.
- Manages **state** (trainable **weights**, non-trainable weights).
- Tracks **losses** and **metrics**.
- Automated compatibility checks (static **shape inference**)

## What does a Layer do?

- Computation from a **batch** of inputs to a **batch** of outputs.
- Manages **state** (trainable **weights**, non-trainable weights).
- Tracks **losses** and **metrics**.
- Automated compatibility checks (static **shape inference**).
- Can be frozen (useful in **fine-tuning** and Transfer Learning).

## What does a Layer do?

- Computation from a **batch** of inputs to a **batch** of outputs.
- Manages **state** (trainable **weights**, non-trainable weights).
- Tracks **losses** and **metrics**.
- Automated compatibility checks (static **shape inference**).
- Can be frozen (useful in **fine-tuning** and Transfer Learning).
- Can be serialized and deserialized (useful for storing the model).

## What does a Layer do?

- Computation from a **batch** of inputs to a **batch** of outputs.
- Manages **state** (trainable **weights**, non-trainable weights).
- Tracks **losses** and **metrics**.
- Automated compatibility checks (static **shape inference**).
- Can be frozen (useful in **fine-tuning** and Transfer Learning).
- Can be serialized and deserialized (useful for storing the model).

# Model Design

- Keras **Sequential** API

- + standard layers

- Stack of layers

- For Simple models



# Model Design

- Keras **Sequential** API

- + standard layers

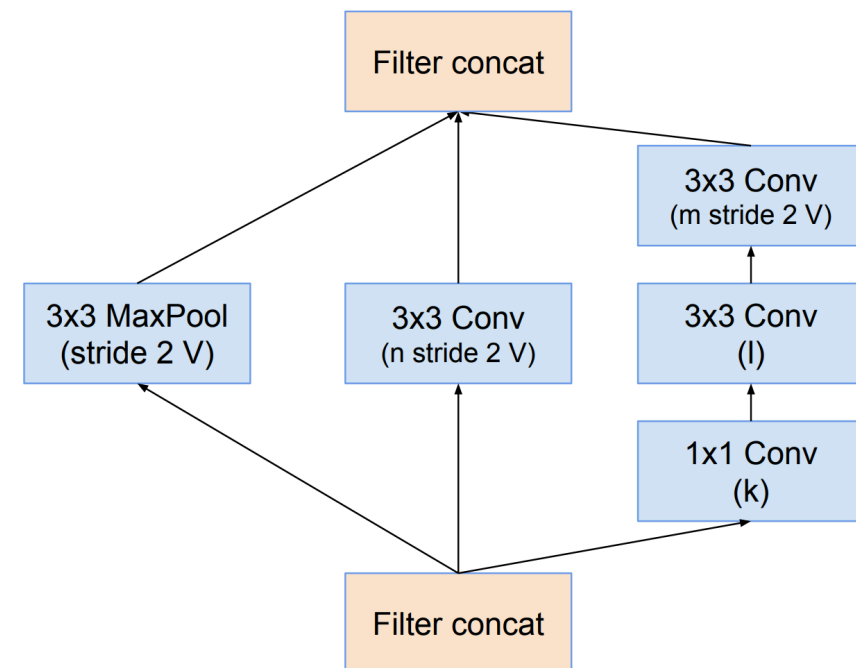
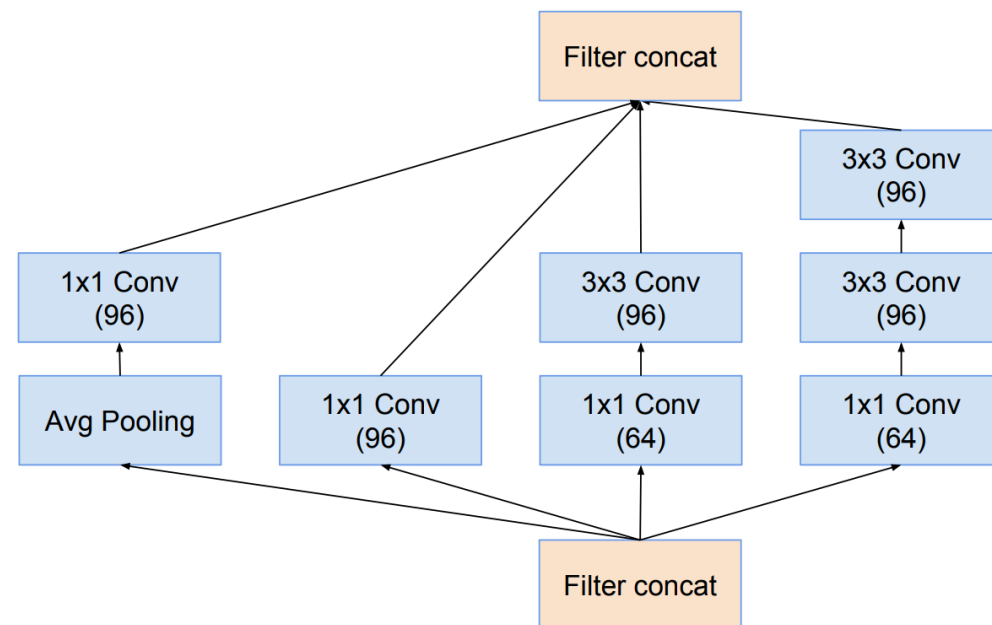
- Stack of layers  
For Simple models

- Keras **Functional** API

- + standard layers

- DAG of layers  
For Simple models

# Functional API (Creating a **DAG**)



## Functional API: A way to define DAGs of layers

```
import tensorflow as tf
from tensorflow.keras import layers

inputs = tf.keras.Input(shape=(784,))

x = layers.Dense(64, activation='relu')(inputs)
x = layers.Dense(64, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = tf.keras.Model(inputs=inputs, outputs=outputs)
```



# Functional API: A way to define DAGs of layers

```
import tensorflow as tf
from tensorflow.keras import layers
```

```
inputs = tf.keras.Input(shape=(784,))
```

```
x = layers.Dense(64, activation='relu')(inputs)
x = layers.Dense(64, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)
```

```
model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

You should first  
specify the model's  
input

## Functional API: A way to define DAGs of layers

```
import tensorflow as tf
from tensorflow.keras import layers
```

```
inputs = tf.keras.Input(shape=(784,))
```

```
x = layers.Dense(64, activation='relu')(inputs)
x = layers.Dense(64, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)
```

Then, define the  
model

```
model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

## Functional API: A way to define DAGs of layers

```
import tensorflow as tf
from tensorflow.keras import layers
```

```
inputs = tf.keras.Input(shape=(784,))
```

```
x = layers.Dense(64, activation='relu')(inputs)
```

```
x = layers.Dense(64, activation='relu')(x)
```

```
outputs = layers.Dense(10, activation='softmax')(x)
```

```
model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

And finally, build  
the model

## Functional API: A way to define DAGs of layers

```
import tensorflow as tf
from tensorflow.keras import layers

inputs = tf.keras.Input(shape=(784,))

x = layers.Dense(64, activation='relu')(inputs)
x = layers.Dense(64, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = tf.keras.Model(inputs=inputs, outputs=outputs)
```

# Functional API: A way to define DAGs of layers

```
model.summary()
```

# Functional API: A way to define DAGs of layers

```
model.summary()
```

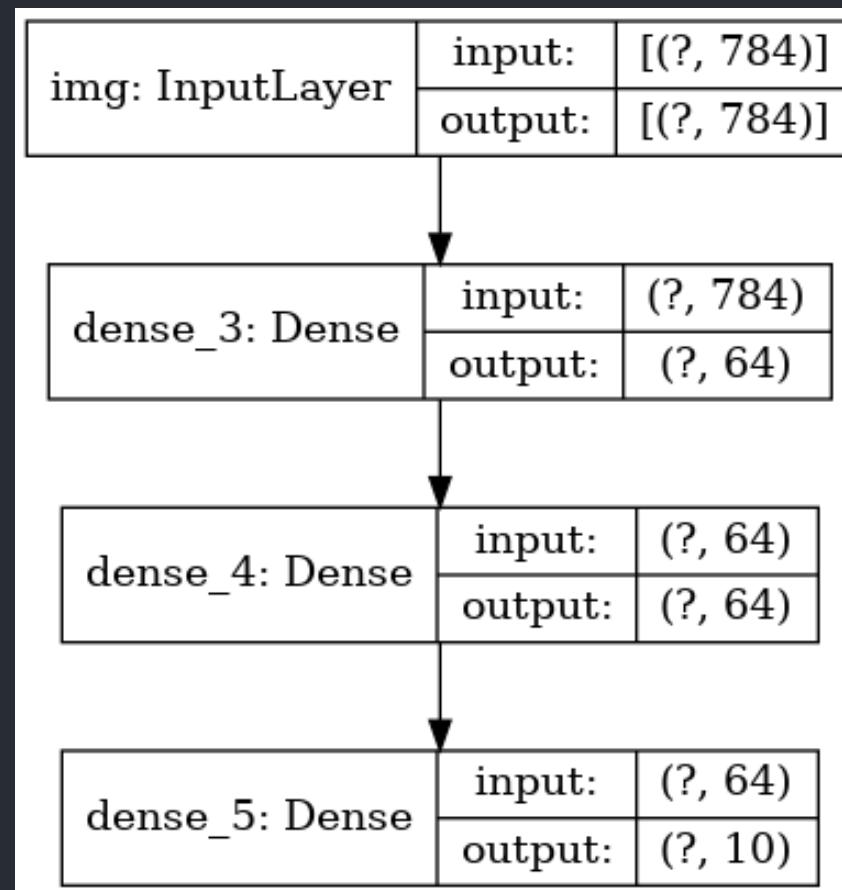
```
-----
Layer (type)                 Output Shape              Param #
=====
img (InputLayer)            [(None, 784)]             0
-----
dense_3 (Dense)              (None, 64)                50240
-----
dense_4 (Dense)              (None, 64)                4160
-----
dense_5 (Dense)              (None, 10)                650
=====
Total params: 55,050
Trainable params: 55,050
Non-trainable params: 0
-----
```

## Functional API: A way to define DAGs of layers

```
keras.utils.plot_model(model, 'plot.png', show_shapes=True)
```

# Functional API: A way to define DAGs of layers

```
keras.utils.plot_model(model, 'plot.png', show_shapes=True)
```





# Example!

## Visual Question Answering

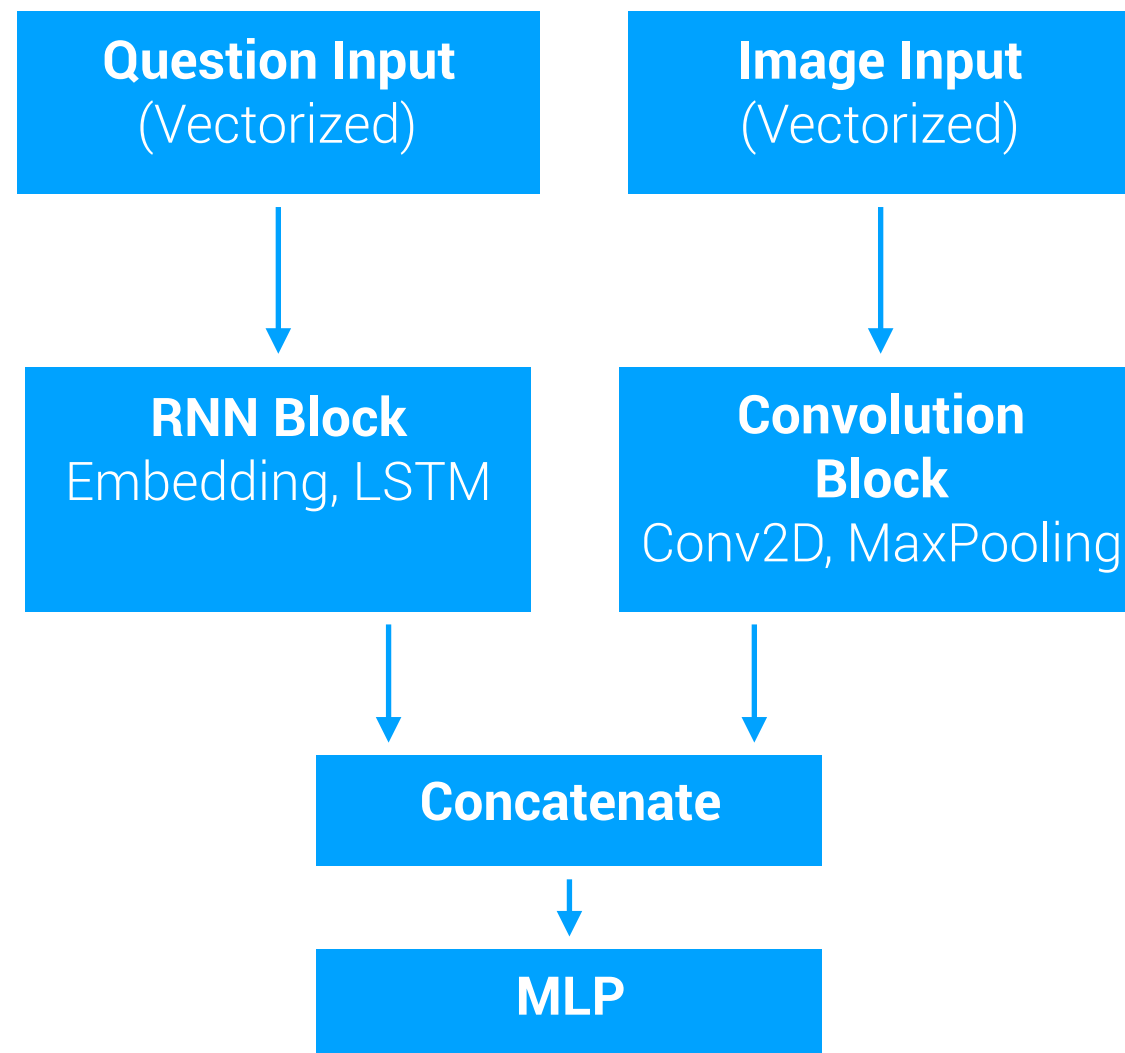


Question:  
**What animal are these?**

Answer:  
**Koala**

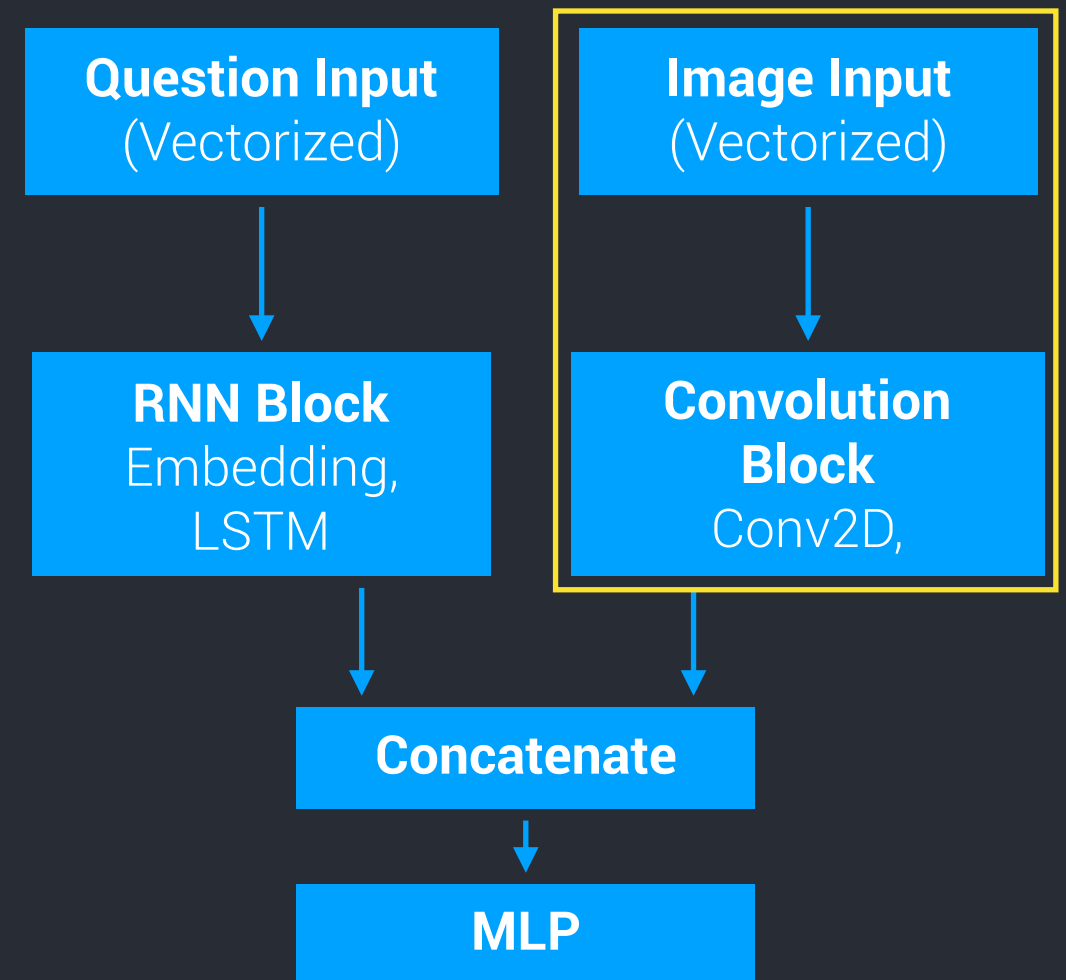
# Example!

## Visual Question Answering



# VQA Example!

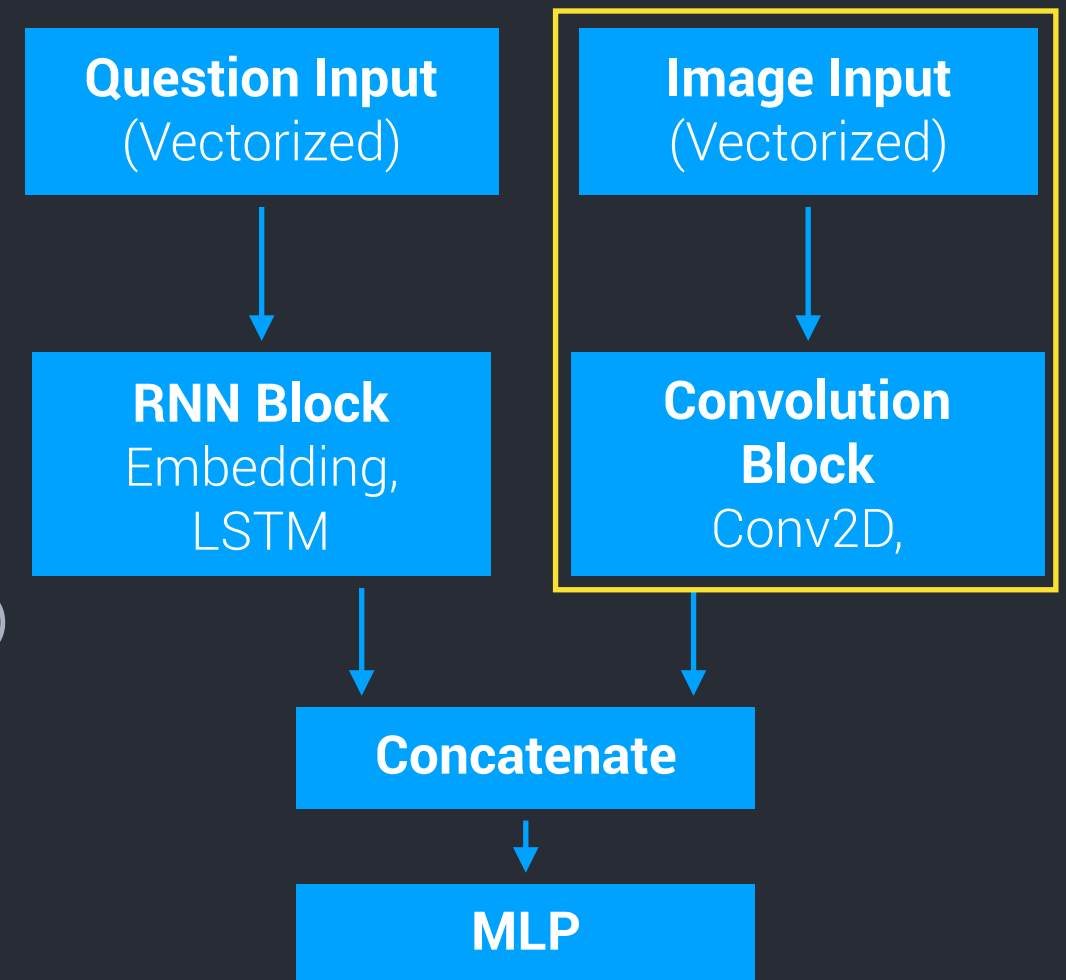
```
# image input  
image_input = Input(shape=(128, 128, 3))
```



# VQA Example!

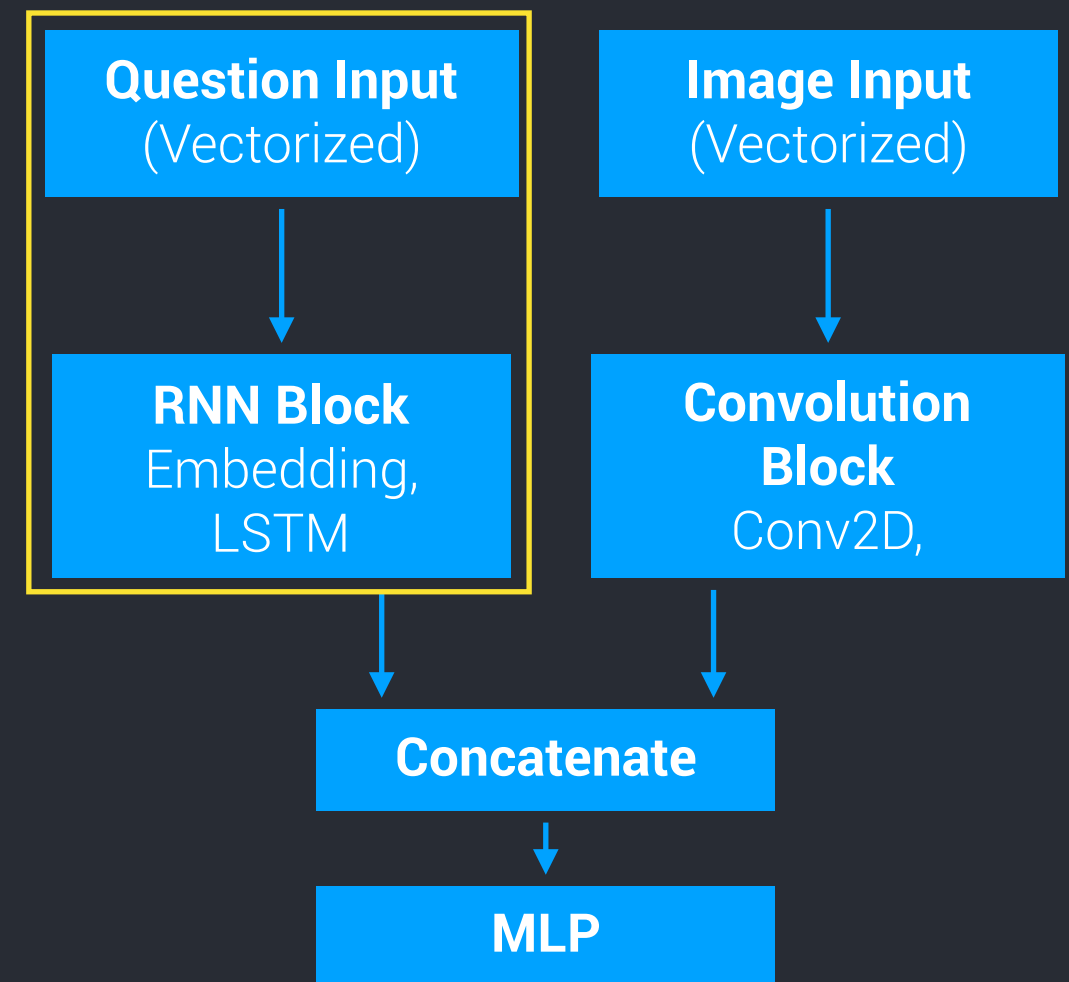
```
# image input
image_input = Input(shape=(128, 128, 3))

# Encode the image into an abstract
# representation
encoded_image = Conv2D(64, (3, 3),
                      activation='relu')(image_input)
encoded_image = MaxPooling2D()(encoded_image)
encoded_image = Flatten()(encoded_image)
```



# VQA Example!

```
# Vectorized input question
question_input = Input(shape=(None,),
                        dtype='int32')
```

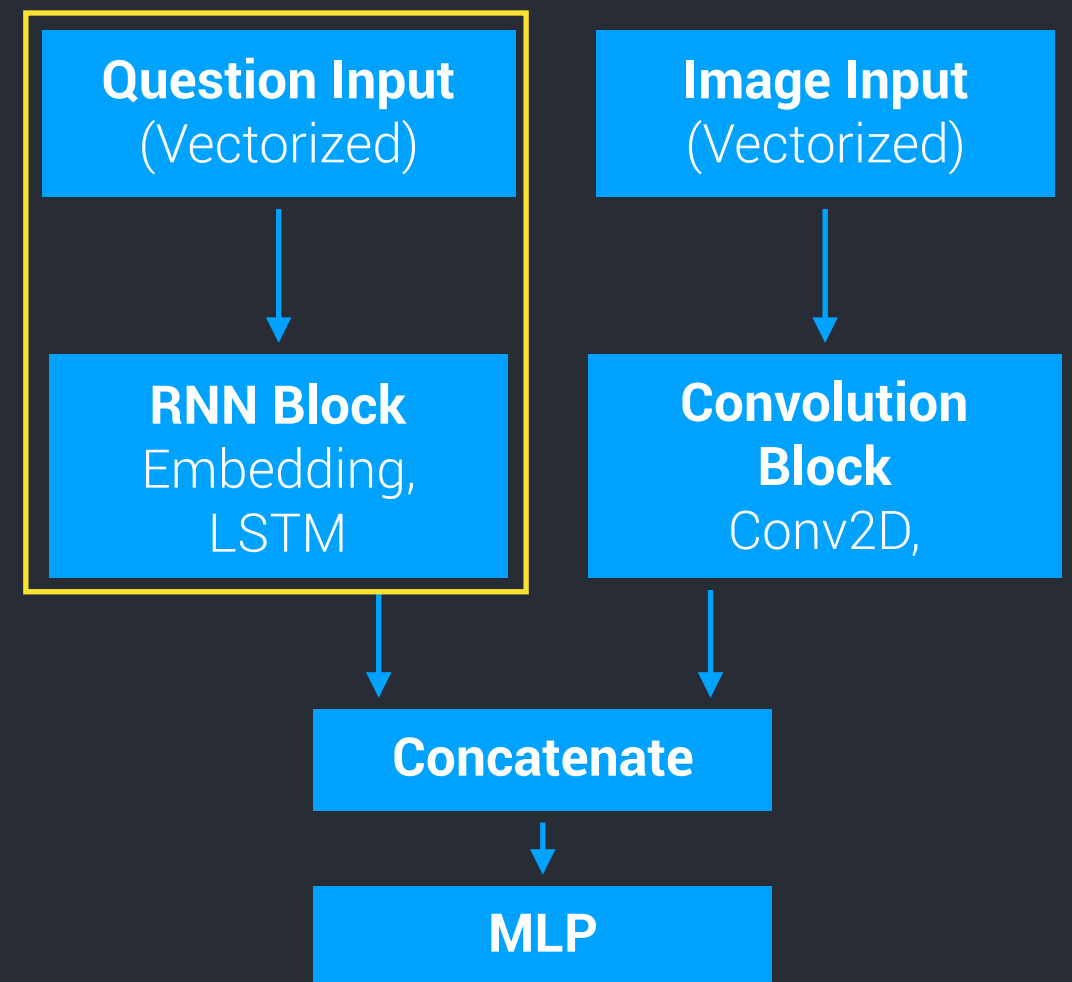


# VQA Example!

```
# Vectorized input question
question_input = Input(shape=(None,),
                        dtype='int32')

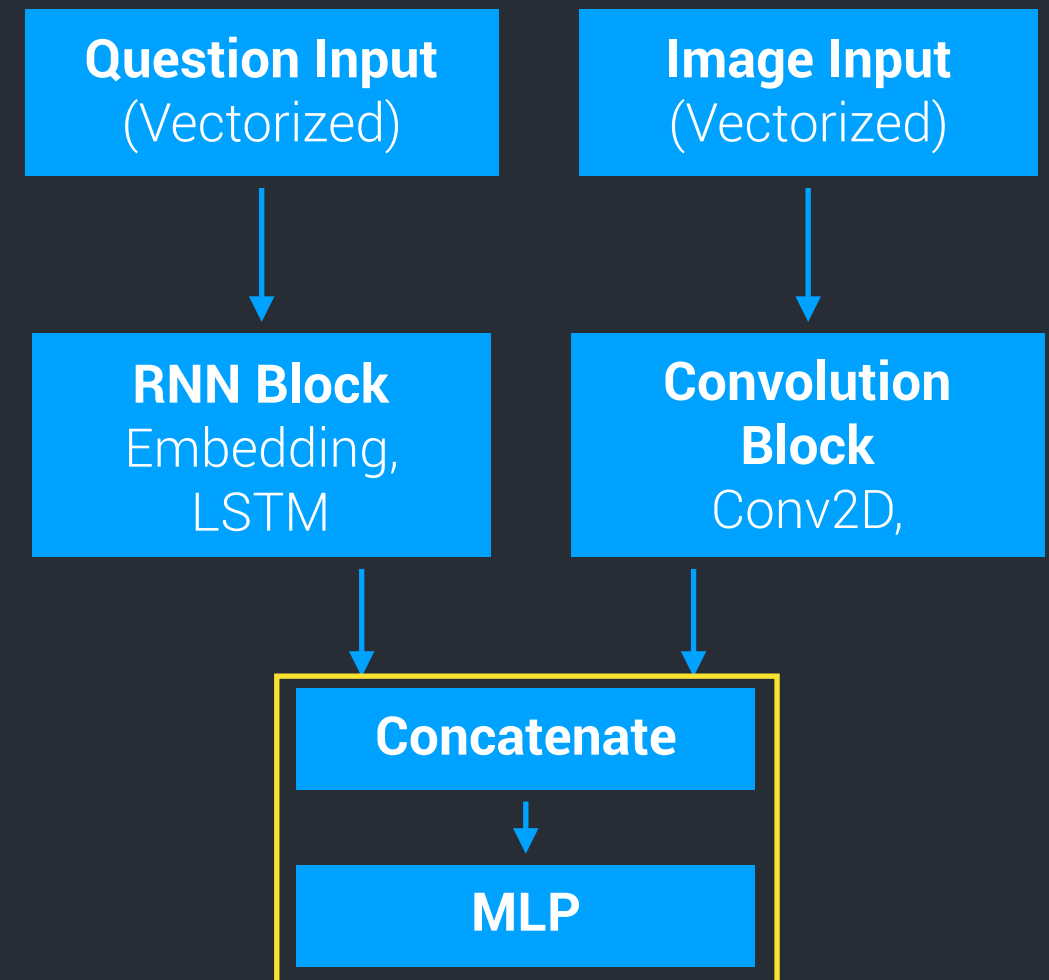
# Encode the question into a single-
# vector representation
embedded = Embedding(
    input_dim=5000,
    output_dim=128,
    mask_zero=True)(question_input)

encoded_question = LSTM(128)(embedded)
```



# VQA Example!

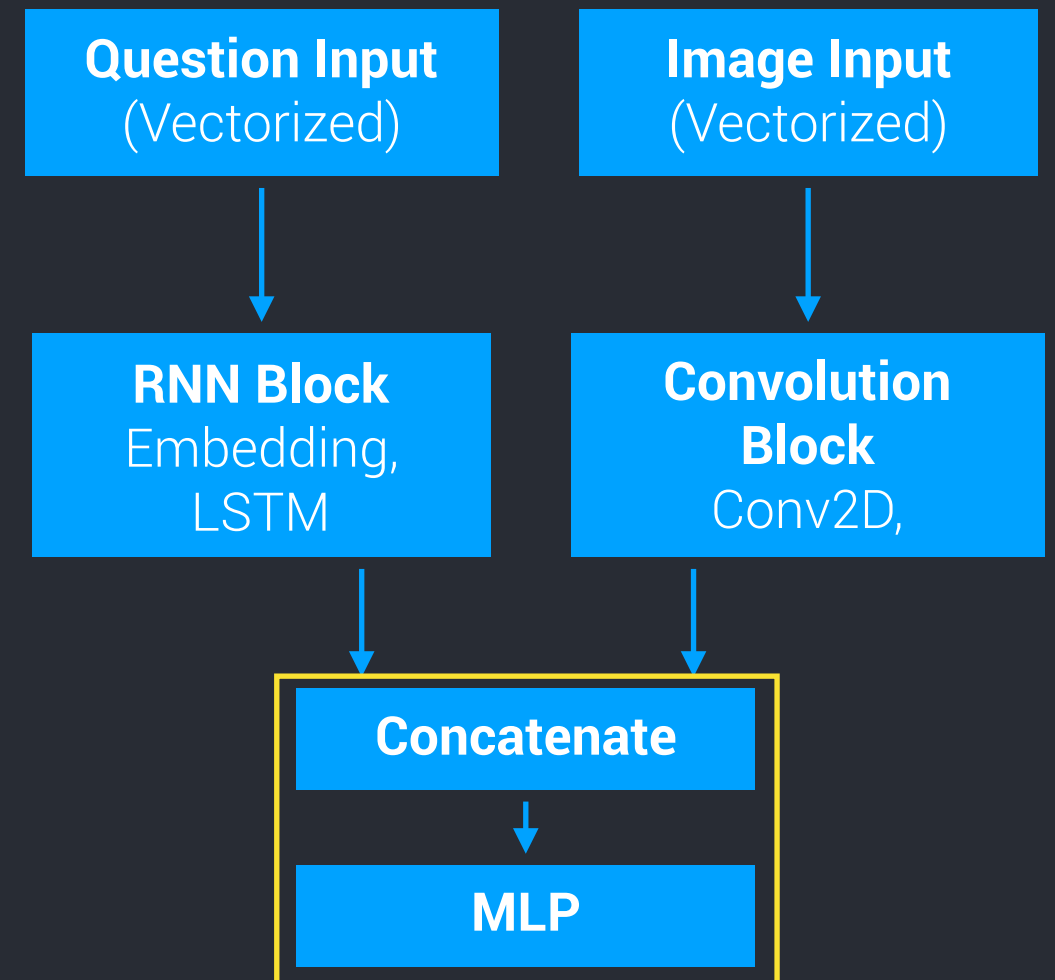
```
# Concat the vector representations  
merged = layers.concatenate([  
    encoded_image, encoded_question])
```



# VQA Example!

```
# Concat the vector representations
merged = layers.concatenate([
    encoded_image, encoded_question])

# Use an MLP to produce the output
output = Dense(1000,
    activation='softmax')(merged)
```





## Quiz #1: Product Review Classifier

Suppose that we have an online store (e.g., Amazon), and users can put a comment on products if they have bought them. Then, we want to find 1) whether the user would like to recommend the product 2) the sentiment of that review. Your model is given the title, the body, and the category of the review.

Here are the details of inputs and outputs:

### Inputs

- **Title:** Vectorized & padded input (can consist of multiple word)
- **Body:** Vectorized & padded review content
- **Product Category:** one category out of 12 (one-hot representation)

### Output

- **Sentiment score:** 5 possibilities
- **Recommend:** Whether the user recommends the product

# Writing Custom Layers

# Custom Layer Outline

```
class MyLayer(layers.Layer):  
    def __init__(self, arg1, arg2, ...):  
        super(MyLayer, self).__init__()  
        ...  
  
    def build(self, input_shape):  
        ...  
  
    def compute_output_shape(self, input_shape):  
        ...  
  
    def compute_mask(self, inputs, mask=None):  
        ...  
  
    def call(self, inputs):  
        ...  
  
    def get_config(self):  
        ...
```

# Custom Layer Outline

```
class MyLayer(layers.Layer):  
    def __init__(self, arg1, arg2, ...):  
        super(MyLayer, self).__init__()  
        ...  
  
    def build(self, input_shape):  
        ...  
  
    def compute_output_shape(self, input_shape):  
        ...  
  
    def compute_mask(self, inputs, mask=None):  
        ...  
  
    def call(self, inputs):  
        ...  
  
    def get_config(self):  
        ...
```

**Required!**

# A Very Basic Layer

# A Very Basic Layer

```
class Linear(layers.Layer):
```

# A Very Basic Layer

```
class Linear(layers.Layer):  
    def __init__(self, units=32, input_dim=32):
```

```
        def call(self, inputs):  
            ...
```

# A Very Basic Layer

```
class Linear(layers.Layer):
    def __init__(self, units=32, input_dim=32):
        super(Linear, self).__init__()

        initializer = tf.initializers.GlorotUniform()
        self.w = tf.Variable(initializer([input_dim, units]),
                             name="kernel")

        initializer = tf.initializers.Zeros()
        self.b = tf.Variable(initializer([units]),
                             name="bias")

    def call(self, inputs):
        ...
```



# A Very Basic Layer

```
class Linear(layers.Layer):
    def __init__(self, units=32, input_dim=32):
        super(Linear, self).__init__()

        initializer = tf.initializers.GlorotUniform()
        self.w = tf.Variable(initializer([input_dim, units]),
                             name="kernel")

        initializer = tf.initializers.Zeros()
        self.b = tf.Variable(initializer([units]),
                             name="bias")

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b
```

# A Very Basic Layer

```
class Linear(layers.Layer):
    def __init__(self, units=32, input_dim=32):
        super(Linear, self).__init__()

        initializer = tf.initializers.GlorotUniform()
        self.w = tf.Variable(initializer([input_dim, units]),
                             name="kernel")

        initializer = tf.initializers.Zeros()
        self.b = tf.Variable(initializer([units]),
                             name="bias")

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b

x = tf.ones((2, 2))
linear_layer = Linear(4, 2)
y = linear_layer(x)
print(y)
```

# A Very Basic Layer

```
class Linear(layers.Layer):
    def __init__(self, units=32, input_dim=32):
        super(Linear, self).__init__()

        initializer = tf.initializers.GlorotUniform()
        self.w = tf.Variable(initializer([input_dim, units]),
                             name="kernel")

        initializer = tf.initializers.Zeros()
        self.b = tf.Variable(initializer([units]),
                             name="bias")
```

**What is the difference between our custom and Keras' Dense layer?**

```
def call(self, inputs):
    return tf.matmul(inputs, self.w) + self.b
```

**Linear(units=..., input\_dim=...)**

```
x = tf.ones((2, 2))
linear_layer = Linear(4, 2)
y = linear_layer(x)
print(y)
```

**Dense(units=...)**

# A Very Basic Layer

```
class Linear(layers.Layer):
    def __init__(self, units=32, input_dim=32):
        super(Linear, self).__init__()

        initializer = tf.initializers.GlorotUniform()
        self.w = tf.Variable(initializer([input_dim, units]),
                             name="kernel")

        initializer = tf.initializers.Zeros()
        self.b = tf.Variable(initializer([units]),
                             name="bias")
```

**What is the difference between our custom and Keras' Dense layer?**

```
def call(self, inputs):
    return tf.matmul(inputs, self.w) + self.b
```

**Linear(units=..., input\_dim=...)**

**Dense(units=...)**

```
x = tf.ones((2, 2))
linear_layer = Linear(4, 2)
y = linear_layer(x)
print(y)
```

Do we have to also specify the input dimension for the Dense layer?

## A Very Basic Layer - v2.0: Adding Laziness!

## A Very Basic Layer - v2.0: Adding Laziness!

```
class Linear(layers.Layer):
    def __init__(self, units=32, input_dim=32):
        super(Linear, self).__init__()

        initializer = tf.initializers.GlorotUniform()
        self.w = tf.Variable(initializer([input_dim, units]),
                             name="kernel")

        initializer = tf.initializers.Zeros()
        self.b = tf.Variable(initializer([units]),
                             name="bias")

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b

x = tf.ones((2, 2))
linear_layer = Linear(4, 2)
y = linear_layer(x)
print(y)
```

## A Very Basic Layer - v2.0: Adding Laziness!

```
class Linear(layers.Layer):  
    def __init__(self, units=32):  
        super(Linear, self).__init__()  
  
    def call(self, inputs):  
        return tf.matmul(inputs, self.w) + self.b
```

```
x = tf.ones((2, 2))  
linear_layer = Linear(4, 2)  
y = linear_layer(x)  
print(y)
```

## A Very Basic Layer - v2.0: Adding Laziness!

```
class Linear(layers.Layer):  
    def __init__(self, units=32):  
        super(Linear, self).__init__()  
        self.units = units  
  
    def call(self, inputs):  
        return tf.matmul(inputs, self.w) + self.b
```

```
x = tf.ones((2, 2))  
linear_layer = Linear(4, 2)  
y = linear_layer(x)  
print(y)
```



## A Very Basic Layer - v2.0: Adding Laziness!

```
class Linear(layers.Layer):
    def __init__(self, units=32):
        super(Linear, self).__init__()
        self.units = units

    def build(self, input_shape):
        self.w = self.add_weight(shape=(input_shape[-1], self.units),
                                initializer='random_normal',
                                trainable=True)
        self.b = self.add_weight(shape=(self.units,),
                                initializer='zeros',
                                trainable=True)

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b

x = tf.ones((2, 2))
linear_layer = Linear(4, 2)
y = linear_layer(x)
print(y)
```

## A Very Basic Layer - v2.0: Adding Laziness!

```
class Linear(layers.Layer):
    def __init__(self, units=32):
        super(Linear, self).__init__()
        self.units = units

    def build(self, input_shape):
        self.w = self.add_weight(shape=(input_shape[-1], self.units),
                                initializer='random_normal',
                                trainable=True)
        self.b = self.add_weight(shape=(self.units,),
                                initializer='zeros',
                                trainable=True)

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b

x = tf.ones((2, 2))
linear_layer = Linear(4)
y = linear_layer(x)
print(y)
```

## A Very Basic Layer - v2.1: Make it Serializable

```
class Linear(layers.Layer):
    def __init__(self, units=32):
        super(Linear, self).__init__()
        self.units = units

    def build(self, input_shape):
        self.w = self.add_weight(shape=(input_shape[-1], self.units),
                                initializer='random_normal',
                                trainable=True)
        self.b = self.add_weight(shape=(self.units,),
                                initializer='zeros',
                                trainable=True)

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b
```

## A Very Basic Layer - v2.1: Make it Serializable

```
class Linear(layers.Layer):
    def __init__(self, units=32):
        super(Linear, self).__init__()
        self.units = units

    def build(self, input_shape):
        self.w = self.add_weight(shape=(input_shape[-1], self.units),
                                initializer='random_normal',
                                trainable=True)
        self.b = self.add_weight(shape=(self.units,),
                                initializer='zeros',
                                trainable=True)

    def call(self, inputs):
        return tf.matmul(inputs, self.w) + self.b

    def get_config(self):
        config = super(Linear, self).get_config()
        config.update({'units': self.units})
        return config
```

## A Very Basic Layer - v2.1: Make it Serializable

```
layer = Linear(64)
config = layer.get_config()
print(config)
```

## A Very Basic Layer - v2.1: Make it Serializable

```
layer = Linear(64)
config = layer.get_config()
print(config)
```

```
{'name': 'linear', 'trainable': True, 'dtype': 'float32', 'units': 64}
```

## A Very Basic Layer - v2.1: Make it Serializable

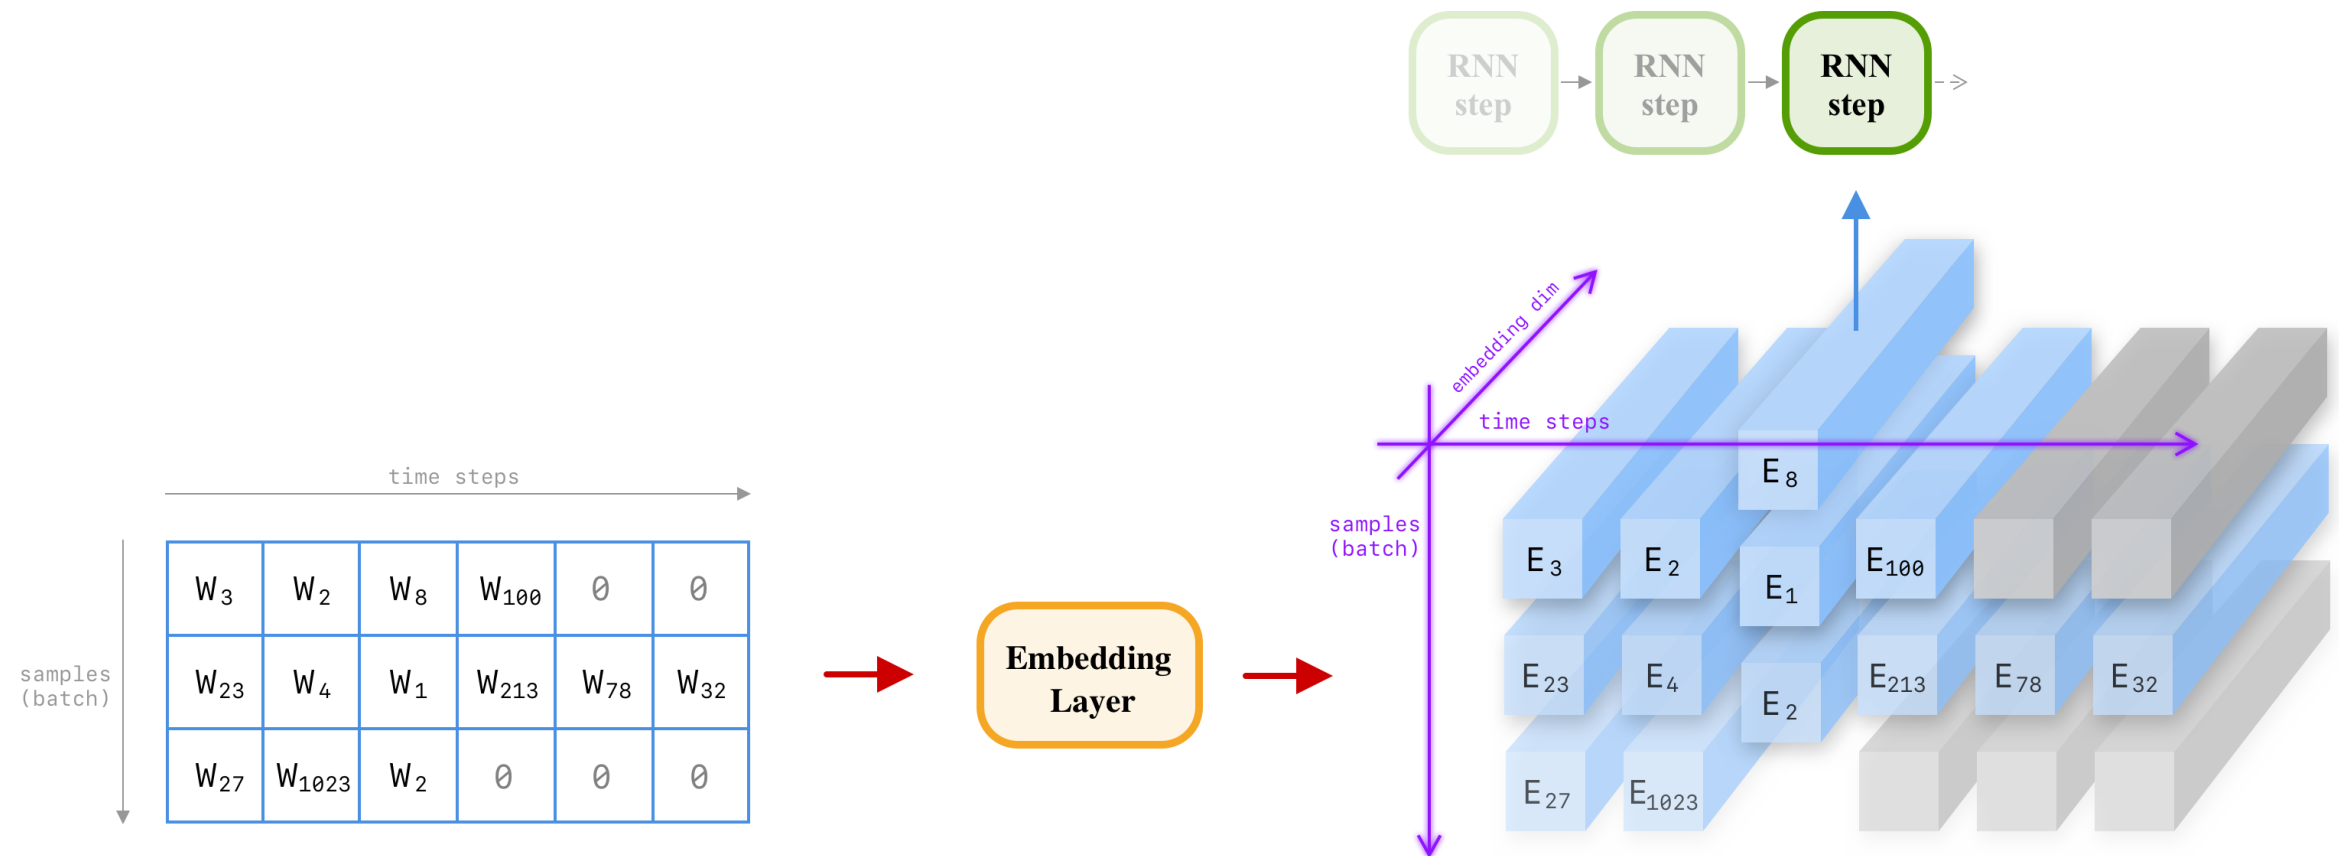
```
layer = Linear(64)
config = layer.get_config()
print(config)

new_layer = Linear.from_config(config)
```

# Masking



# Masking



# Masking

Keras Layers fall into  
**3 categories** when it  
comes to making:

1. Mask **Consumers**
2. Mask **Propagators**
3. Mask **Generators**

# Masking in Keras

```
class ConsumerLayer(layers.Layer):  
    def call(self, inputs):  
        ...
```

# Masking in Keras

```
class ConsumerLayer(layers.Layer):  
    def call(self, inputs, mask=None):  
        ...
```

# Masking in Keras

```
class ConsumerLayer(layers.Layer):  
    def call(self, inputs, mask=None):  
        ...
```

```
class MaskPassThroughLayer(layers.Layer):  
    def __init__(self, ...):  
        self.support_masking = True
```

# Masking in Keras

```
class ConsumerLayer(layers.Layer):  
    def call(self, inputs, mask=None):  
        ...
```

```
class MaskPassThroughLayer(layers.Layer):  
    def __init__(self, ...):  
        self.support_masking = True
```

```
class GeneratorLayer(layers.Layer):  
    def __init__(self, ...):  
        self.support_masking = True  
  
    def compute_mask(self, inputs, mask=None):  
        ...
```

## Masking in Keras: Example

```
class CustomEmbedding(tf.keras.layers.Layer):  
    def __init__(self, input_dim, output_dim, mask_zero=False):  
        super(CustomEmbedding, self).__init__()  
        self.mask_zero = mask_zero  
  
    def build(self, input_shape):  
        ...  
    def call(self, inputs):  
        ...
```

## Masking in Keras: Example

```
class CustomEmbedding(tf.keras.layers.Layer):
    def __init__(self, input_dim, output_dim, mask_zero=False):
        super(CustomEmbedding, self).__init__()
        self.supports_masking = True
        self.mask_zero = mask_zero

    def build(self, input_shape):
        ...
    def call(self, inputs):
        ...
```



# Masking in Keras: Example

```
class CustomEmbedding(tf.keras.layers.Layer):
    def __init__(self, input_dim, output_dim, mask_zero=False):
        super(CustomEmbedding, self).__init__()
        self.supports_masking = True
        self.mask_zero = mask_zero

    def build(self, input_shape):
        ...

    def call(self, inputs):
        ...

    def compute_mask(self, inputs, mask=None):
        if not self.mask_zero:
            return None
        return tf.not_equal(inputs, 0)
```

## Masking in Keras: Example

```
layer = CustomEmbedding(10, 32, mask_zero=True)
x = np.array(
    [[2, 3, 4, 0, 0],
     [3, 3, 4, 9, 20],
     [9, 11, 1, 0, 0]], dtype=np.int32)

y = layer(x)
mask = layer.compute_mask(x)
```

## Masking in Keras: Example

```
layer = CustomEmbedding(10, 32, mask_zero=True)
```

```
x = np.array(  
    [[2, 3, 4, 0, 0],  
     [3, 3, 4, 9, 20],  
     [9, 11, 1, 0, 0]], dtype=np.int32)
```

```
y = layer(x)
```

```
mask = layer.compute_mask(x)
```

```
tf.Tensor(  
[[ True  True  True False False]  
 [ True  True  True  True  True]  
 [ True  True  True False False]], shape=(3, 5), dtype=bool)
```

## training argument in the call method

```
class CustomDropout(layers.Layer):  
    def __init__(self, rate, **kwargs):  
        super(CustomDropout, self).__init__(**kwargs)  
        self.rate = rate  
  
    def call(self, inputs, training=None):  
        ...
```

## training argument in the call method

```
class CustomDropout(layers.Layer):  
    def __init__(self, rate, **kwargs):  
        super(CustomDropout, self).__init__(**kwargs)  
        self.rate = rate  
  
    def call(self, inputs, training=None):  
        if training:  
            ...
```

## Nested Layers

```
class CustomDropout(layers.Layer):  
    def __init__(self, rate):  
        super(CustomDropout, self).__init__()  
        self.rate = rate  
  
    def call(self, inputs, training=None):  
        if training:  
            return tf.nn.dropout(inputs, rate=self.rate)  
        return inputs
```

# Nested Layers

```
class CustomDropout(layers.Layer):  
    def __init__(self, rate):  
        super(CustomDropout, self).__init__()  
        self.rate = rate  
  
    def call(self, inputs, training=None):  
        if training:  
            return tf.nn.dropout(inputs, rate=self.rate)  
        return inputs
```

```
mlp = MLPBlock()  
y = mlp(tf.ones(shape=(3, 64)))  
  
print('trainable weights:', len(mlp.trainable_weights))  
# ?
```

## Nested Layers

```
class CustomDropout(layers.Layer):
    def __init__(self, rate):
        super(CustomDropout, self).__init__()
        self.rate = rate

    def call(self, inputs, training=None):
        if training:
            return tf.nn.dropout(inputs, rate=self.rate)
        return inputs

mlp = MLPBlock()
y = mlp(tf.ones(shape=(3, 64)))

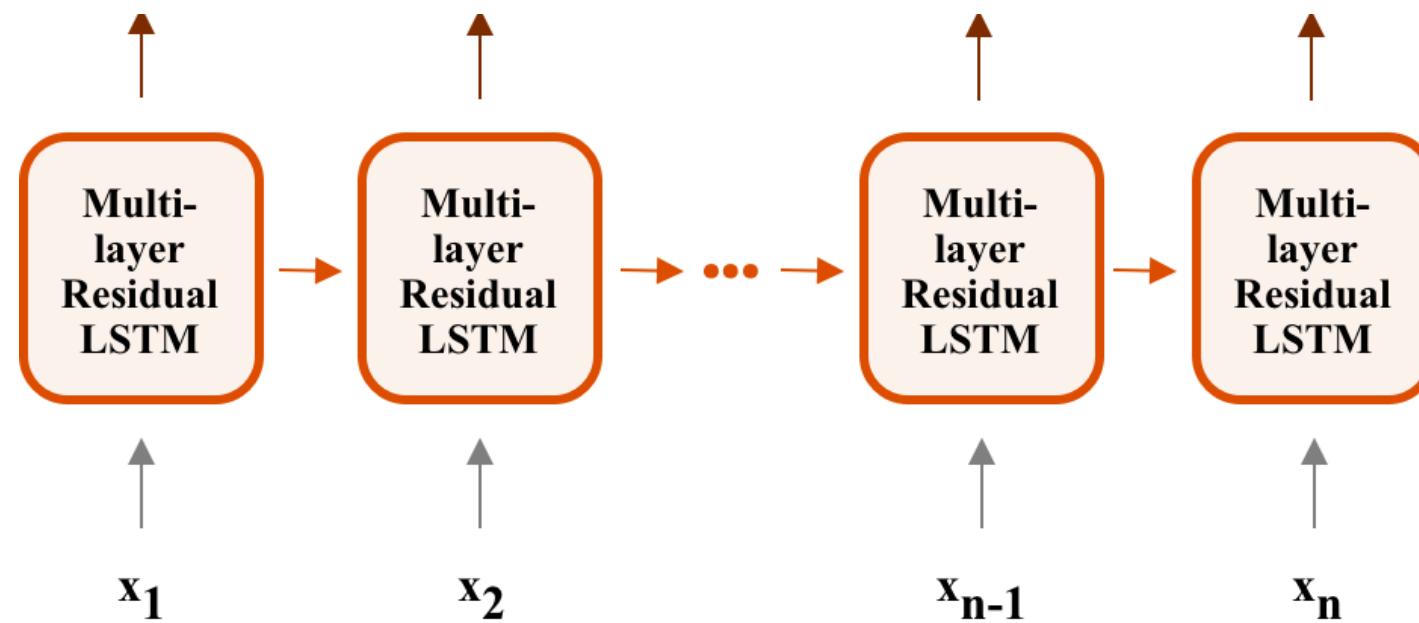
print('trainable weights:', len(mlp.trainable_weights))
# 6
```



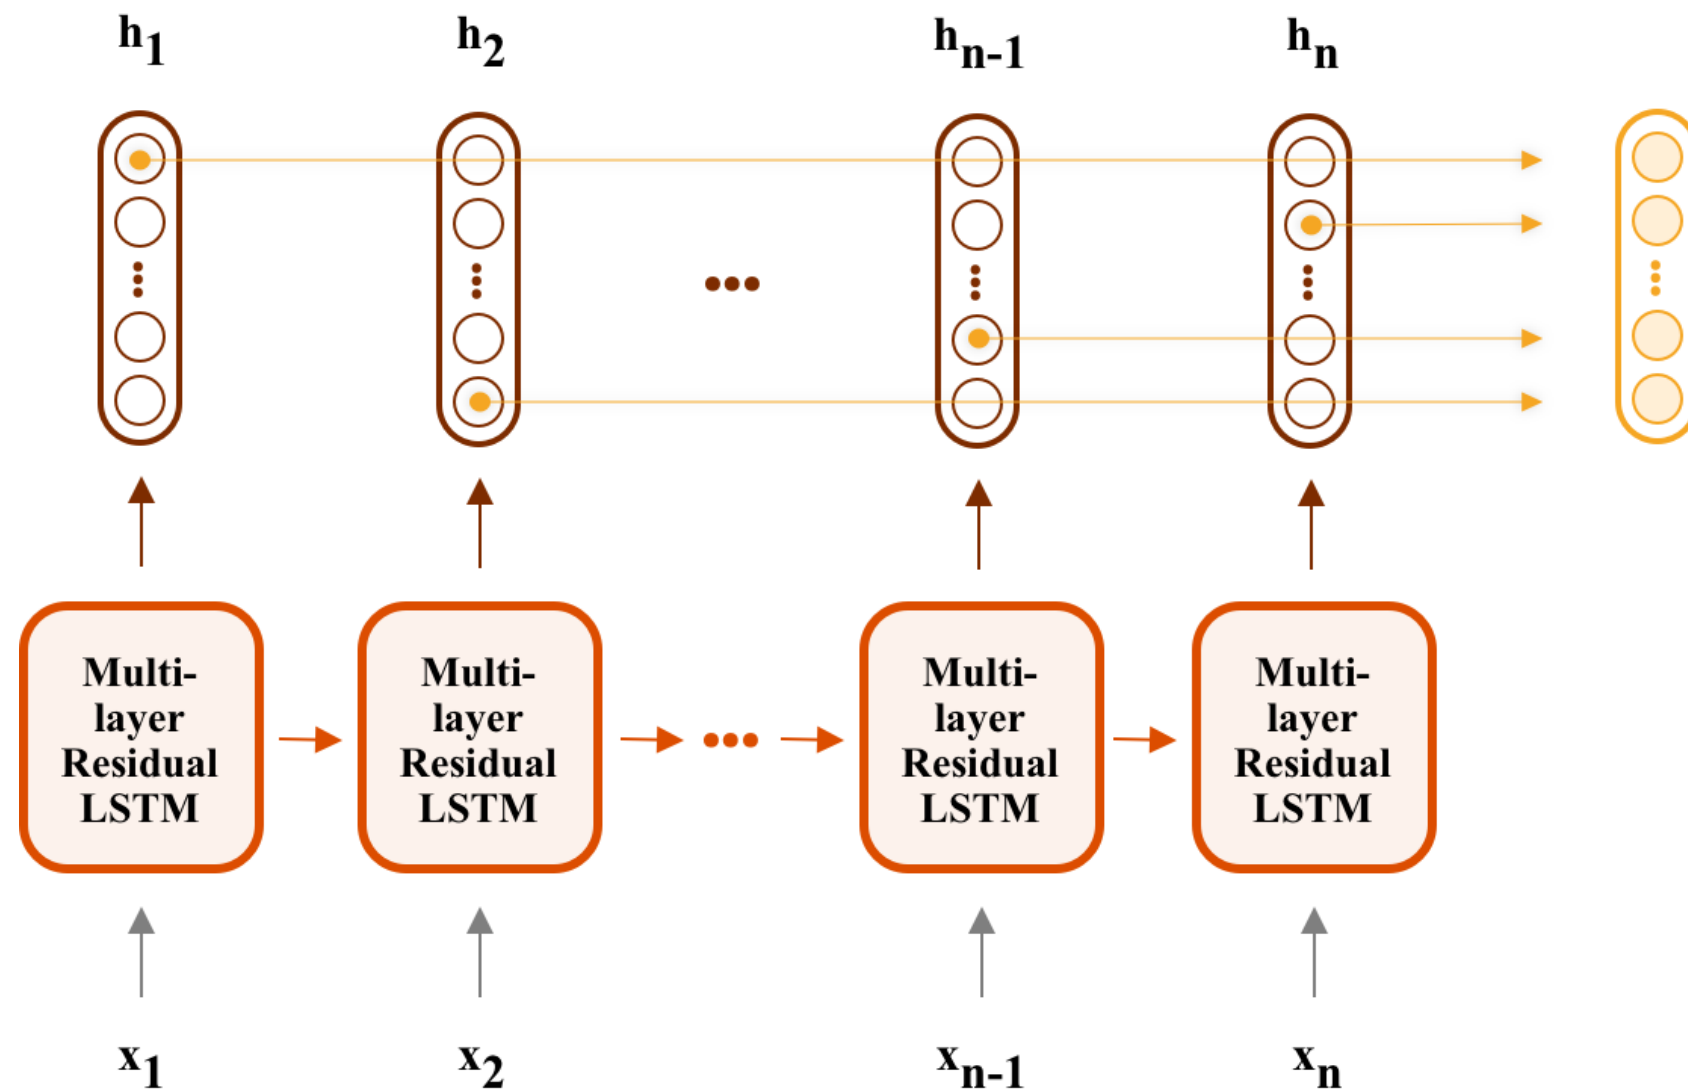
## Quiz #2: Max-pooling through time

One technique that originates from Computer Vision is called Max pooling. As you might remember, this technique reduces the impact of spatial information in the image. For example, If your CNN says, "Yay! I found a wheel at the position (x,y).", your Max-pooling will convert this sentence to "Yay! I found a wheel in this image." Intuitively, we can use the max-pooling procedure in any configuration beside an image. Here is an example of Max-pooling application in recurrent networks:

## Quiz #2: Max-pooling through time



## Quiz #2: Max-pooling through time



## Quiz #2: Max-pooling through time

In this setup, we'd like to perform max-pooling over the hidden states  $\hat{h} = \text{MaxPool}([h^{(1)}, \dots, h^{(n)}])$  where  $h$  is the max-pooled version. Every dim of  $h$  is the maximum of that particular dim across all of the hidden states.

$$\hat{h}_i = \max_{1 \leq k \leq n} h_i^{(k)}$$

Although the default Keras framework provides the implementation, it lacks the masking support. Implement this mechanism as a Keras layer.

# Model Design

- Keras **Sequential** API

- + standard layers

- Stack of layers  
For Simple models

- Keras **Functional** API

- + standard layers

- DAG of layers  
For Simple models

# Model Design

## ● Keras **Sequential** API

+ standard layers

Stack of layers  
For Simple models

## ● Model **Sub-classing**

+ standard layers

Define model by Python  
For very customized models

## ● Keras **Functional** API

+ standard layers

DAG of layers  
For Simple models

# Custom Model

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__()  
        ...  
  
    def call(self, inputs):  
        # Define your forward pass here
```

# Custom Model

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__()  
        self.dense_1 = layers.Dense(32, activation='relu')  
        self.dense_2 = layers.Dense(num_classes,  
                                     activation='softmax')  
  
    def call(self, inputs):  
        # Define your forward pass here  
        x = self.dense_1(inputs)  
        return self.dense_2(x)
```



# Custom Model

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__()  
        self.dense_1 = layers.Dense(32, activation='relu')  
        self.dense_2 = layers.Dense(num_classes,  
                                     activation='softmax')  
  
    def call(self, inputs):  
        # Define your forward pass here  
        x = self.dense_1(inputs)  
        return self.dense_2(x)
```

# Custom Model

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__()  
        self.dense_1 = layers.Dense(32)  
        self.dense_2 = layers.Dense(num_classes,  
                                     activation='softmax')  
  
    def call(self, inputs):  
        # Define your forward pass here  
        x = self.dense_1(inputs)  
        x = tf.nn.relu(x)  
        return self.dense_2(x)
```

# Models vs. Layers

## Models vs. Layers

**Models are exactly the same as layers!**

## Models vs. Layers

Models are exactly the same as layers! **plus:**

## Models vs. Layers

Models are exactly the same as layers! **plus:**

- + Training (`model.fit()`, `.compile()`, `.evaluate`, and etc.)
- + Save and load on the disk
- + Summary/Visualization

## Models vs. Layers

### Layer

corresponds to what we refer to in the literature as a "layer" (as in "**convolution layer**" or "**recurrent layer**") or as a "block" (as in "**ResNet block**" or "**Inception block**").

### Model

corresponds to what is referred to in the literature as a "model" (as in "**deep learning model**") or as a "network" (as in "**deep neural network**")

# Model Sub-Classing (MSC) vs. Functional API



## Model Sub-Classing (MSC) vs. Functional API

- ✓ **MSC is much more flexible in the graph definition**  
(recall that in Functional API everything should be an instance of a Layer when we are connecting different nodes of the model. Hence, no support for low-level TF Ops)

## Model Sub-Classing (MSC) vs. Functional API

- ✓ MSC is much more flexible in the graph definition  
(recall that in Functional API everything should be an instance of a Layer when we are connecting different nodes of the model. Hence, no support for low-level TF Ops)
- ✓ MSC supports changing the runtime branch between training and evaluation (via training parameter)

## Model Sub-Classing (MSC) vs. Functional API

- ✓ MSC is much more flexible in the graph definition  
(recall that in Functional API everything should be an instance of a Layer when we are connecting different nodes of the model. Hence, no support for low-level TF Ops)
- ✓ MSC supports changing the runtime branch between training and evaluation (via training parameter)
- ✓ The mask argument should be passed manually in the MSC.

## Model Sub-Classing (MSC) vs. Functional API

- ✓ MSC is much more flexible in the graph definition (recall that in Functional API everything should be an instance of a Layer when we are connecting different nodes of the model. Hence, no support for low-level TF Ops)
- ✓ MSC supports changing the runtime branch between training and evaluation (via training parameter)
- ✓ The mask argument should be passed manually in the MSC.
- ✓ Model saving & restoring is easier in the Functional API.

# Model Design

## ● Keras **Sequential** API

+ standard layers

Stack of layers  
For Simple models

## ● Model **Sub-classing**

+ standard layers

Define model by Python  
For very customized models

## Keras **Functional** API

+ standard layers

DAG of layers  
For Simple models

# Model Design

## ● Keras **Sequential** API

- + standard layers
- + custom layers, losses, and metrics
- Stack of layers
- For Simple models
- 
- 

## ● Model **Sub-classing**

- + standard layers
- + custom layers, losses, and metrics
- Define model by Python
- For very customized models
- 
- 

## ■ Keras **Functional** API

- + standard layers
- + custom layers, losses, and metrics
- DAG of layers
- For Simple models



# Model Training



# Model Training

- Keras **built-in** loops

- model.fit()  
+ callbacks

- Fast prototyping





# Keras built-in training loops

```
model = MyModel()
```

## Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

## Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

## Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

- `.compile()` is about configuring the training process.
- Specify the optimizer, loss, and metrics.

## Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

```
history = model.fit(data,  
                    epochs=10, batch_size=128,  
                    validation_data=val_data,  
                    callbacks=[EarlyStopping(),  
                               TensorBoard(),  
                               ModelCheckpoint()])
```

## Keras built-in training loops

```
model = MyModel()

model.compile(optimizer=Adam(),
              loss=BinaryCrossentropy(),
              metrics=[AUC(), Precision(), Recall()])

history = model.fit(data,
                    epochs=10, batch_size=128,
                    validation_data=val_data,
                    callbacks=[EarlyStopping(),
                             TensorBoard(),
                             ModelCheckpoint()])

results = model.evaluate(test_data, batch_size=128)
```

## Keras built-in training loops (run in dynamic graph mode)

```
model = MyModel()

model.compile(optimizer=Adam(),
              loss=BinaryCrossentropy(),
              metrics=[AUC(), Precision(), Recall()],
              run_eagerly=True)

history = model.fit(data,
                    epochs=10, batch_size=128,
                    validation_data=val_data,
                    callbacks=[EarlyStopping(),
                             TensorBoard(),
                             ModelCheckpoint()])

results = model.evaluate(test_data, batch_size=128)

predictions = model.predict(x_test[:3])
```

# Model Training

- Keras **built-in** loops

- model.fit()  
+ callbacks

- Fast prototyping





# Model Training

- Keras **built-in** loops

- model.fit()  
+ callbacks

- Fast prototyping

- **GradientTape**

- custom training loops

- Complete control



## Gradient tapes

Tensorflow "records" all operations executed inside the context of a *tf.GradientTape* onto a "tape".

Tensorflow then uses that tape and the gradients associated with each recorded operation to compute the gradients of a "recorded" computation using reverse mode differentiation.

## Gradient Tape

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    y = x * x

dy_dx = g.gradient(y, x) # Will compute to 6.0
```

# Gradient Tape

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    y = x * x

dy_dx = g.gradient(y, x) # Will compute to 6.0
```

How to calculate second-order derivatives?

## Gradient Tape

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    with tf.GradientTape() as gg:
        gg.watch(x)
        y = x * x
    dy_dx = gg.gradient(y, x)      # Will compute to 6.0

d2y_dx2 = g.gradient(dy_dx, x)   # Will compute to 2.0
```

# Gradient Tape; Real-world usage

# Gradient Tape; Real-world usage

```
lr = 0.05  
m = tf.Variable(0.)  
b = tf.Variable(0.)
```

# Gradient Tape; Real-world usage

```
lr = 0.05  
m = tf.Variable(0.)  
b = tf.Variable(0.)  
  
# Training loop  
for x, y in dataset:
```



# Gradient Tape; Real-world usage

```
lr = 0.05
m = tf.Variable(0.)
b = tf.Variable(0.)

# Training loop
for x, y in dataset:
    with tf.GradientTape() as g:
        preds = model(x)
        loss = loss_fn(y, preds)
```

## Gradient Tape; Real-world usage

```
lr = 0.05
m = tf.Variable(0.)
b = tf.Variable(0.)

# Training loop
for x, y in dataset:
    with tf.GradientTape() as g:
        preds = model(x)
        loss = loss_fn(y, preds)

    g_m, g_b = g.gradient(loss, [m, b])
```

# Gradient Tape; Real-world usage

```
lr = 0.05
m = tf.Variable(0.)
b = tf.Variable(0.)

# Training loop
for x, y in dataset:
    with tf.GradientTape() as g:
        preds = model(x)
        loss = loss_fn(y, preds)

    g_m, g_b = g.gradient(loss, [m, b])

    m.assign_sub(g_m * lr)
    b.assign_sub(g_b * lr)
```

## Gradient Tape; Real-world usage

```
lr = 0.05
m = tf.Variable(0.)
b = tf.Variable(0.)

# Training loop
for x, y in dataset:
    with tf.GradientTape() as g:
        preds = model(x)
        loss = loss_fn(y, preds)

    grads = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(
        zip(grads, model.trainable_variables))
```

## Gradient Tape; Real-world usage

```
lr = 0.05
m = tf.Variable(0.)
b = tf.Variable(0.)

def train_step(x, y):
    with tf.GradientTape() as g:
        preds = model(x)
        loss = loss_fn(y, preds)

    grads = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(
        zip(grads, model.trainable_variables))

    return loss.numpy()
```

## Gradient Tape; Real-world usage

```
lr = 0.05
m = tf.Variable(0.)
b = tf.Variable(0.)

@tf.function
def train_step(x, y):
    with tf.GradientTape() as g:
        preds = model(x)
        loss = loss_fn(y, preds)

    grads = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(
        zip(grads, model.trainable_variables))

    return loss.numpy()
```

## Assignment: Low-level MNIST CNN-Classfier

We have implemented an MNIST classifier several times. However, we want to implement a multi-layer CNN classifier using raw TensorFlow operations and matrix multiplication. No pre-defined CNN blocks are allowed. Moreover, you should organize your model using TF 2.0 Model sub-classing APIs (your assignment will be evaluated based on the cleanness and your effective usages of high-level Model APIs). Additionally, the only available training option is the TF's GradientTape.

Assignments criteria:

- ▶ Matrix implementation of a convolutional layer.
- ▶ Model organization using TF 2.0 sub-classing API
- ▶ Model training using TF 2.0 GradientTape

# Thank you!

