

TensorFlow 2.0 Tutorial: Part #4

High-level APIs for Model training
(`model.fit(..)` & `GradientTape`)



Iran University of Science and Technology (IUST)
Department of Computer Engineering

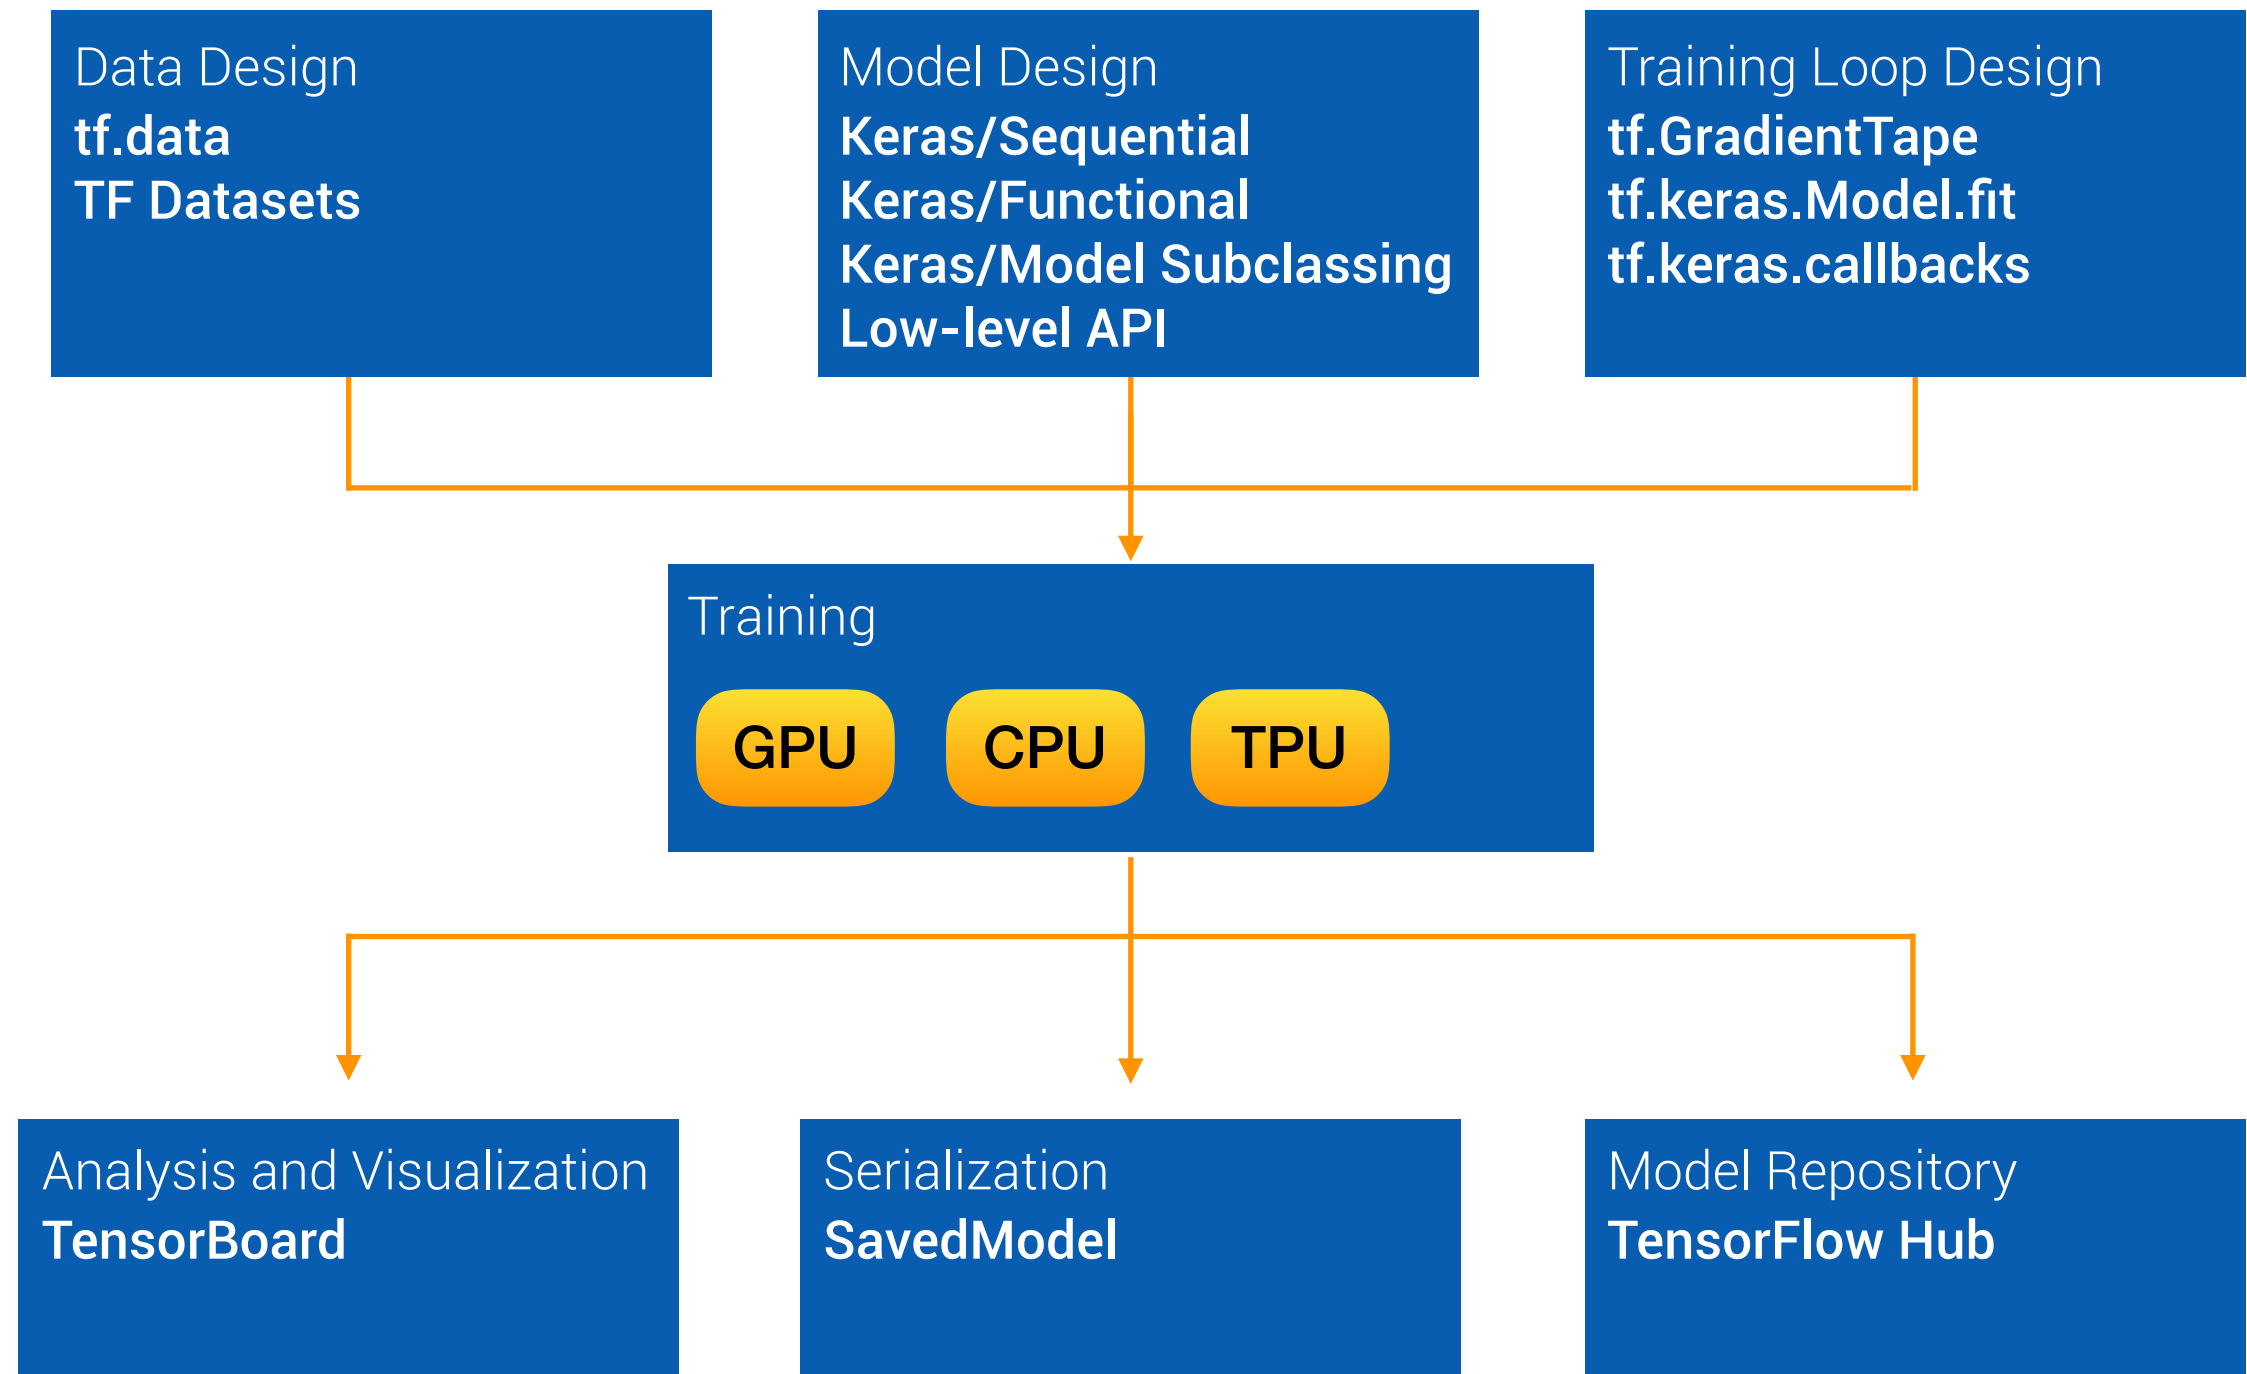


URL:

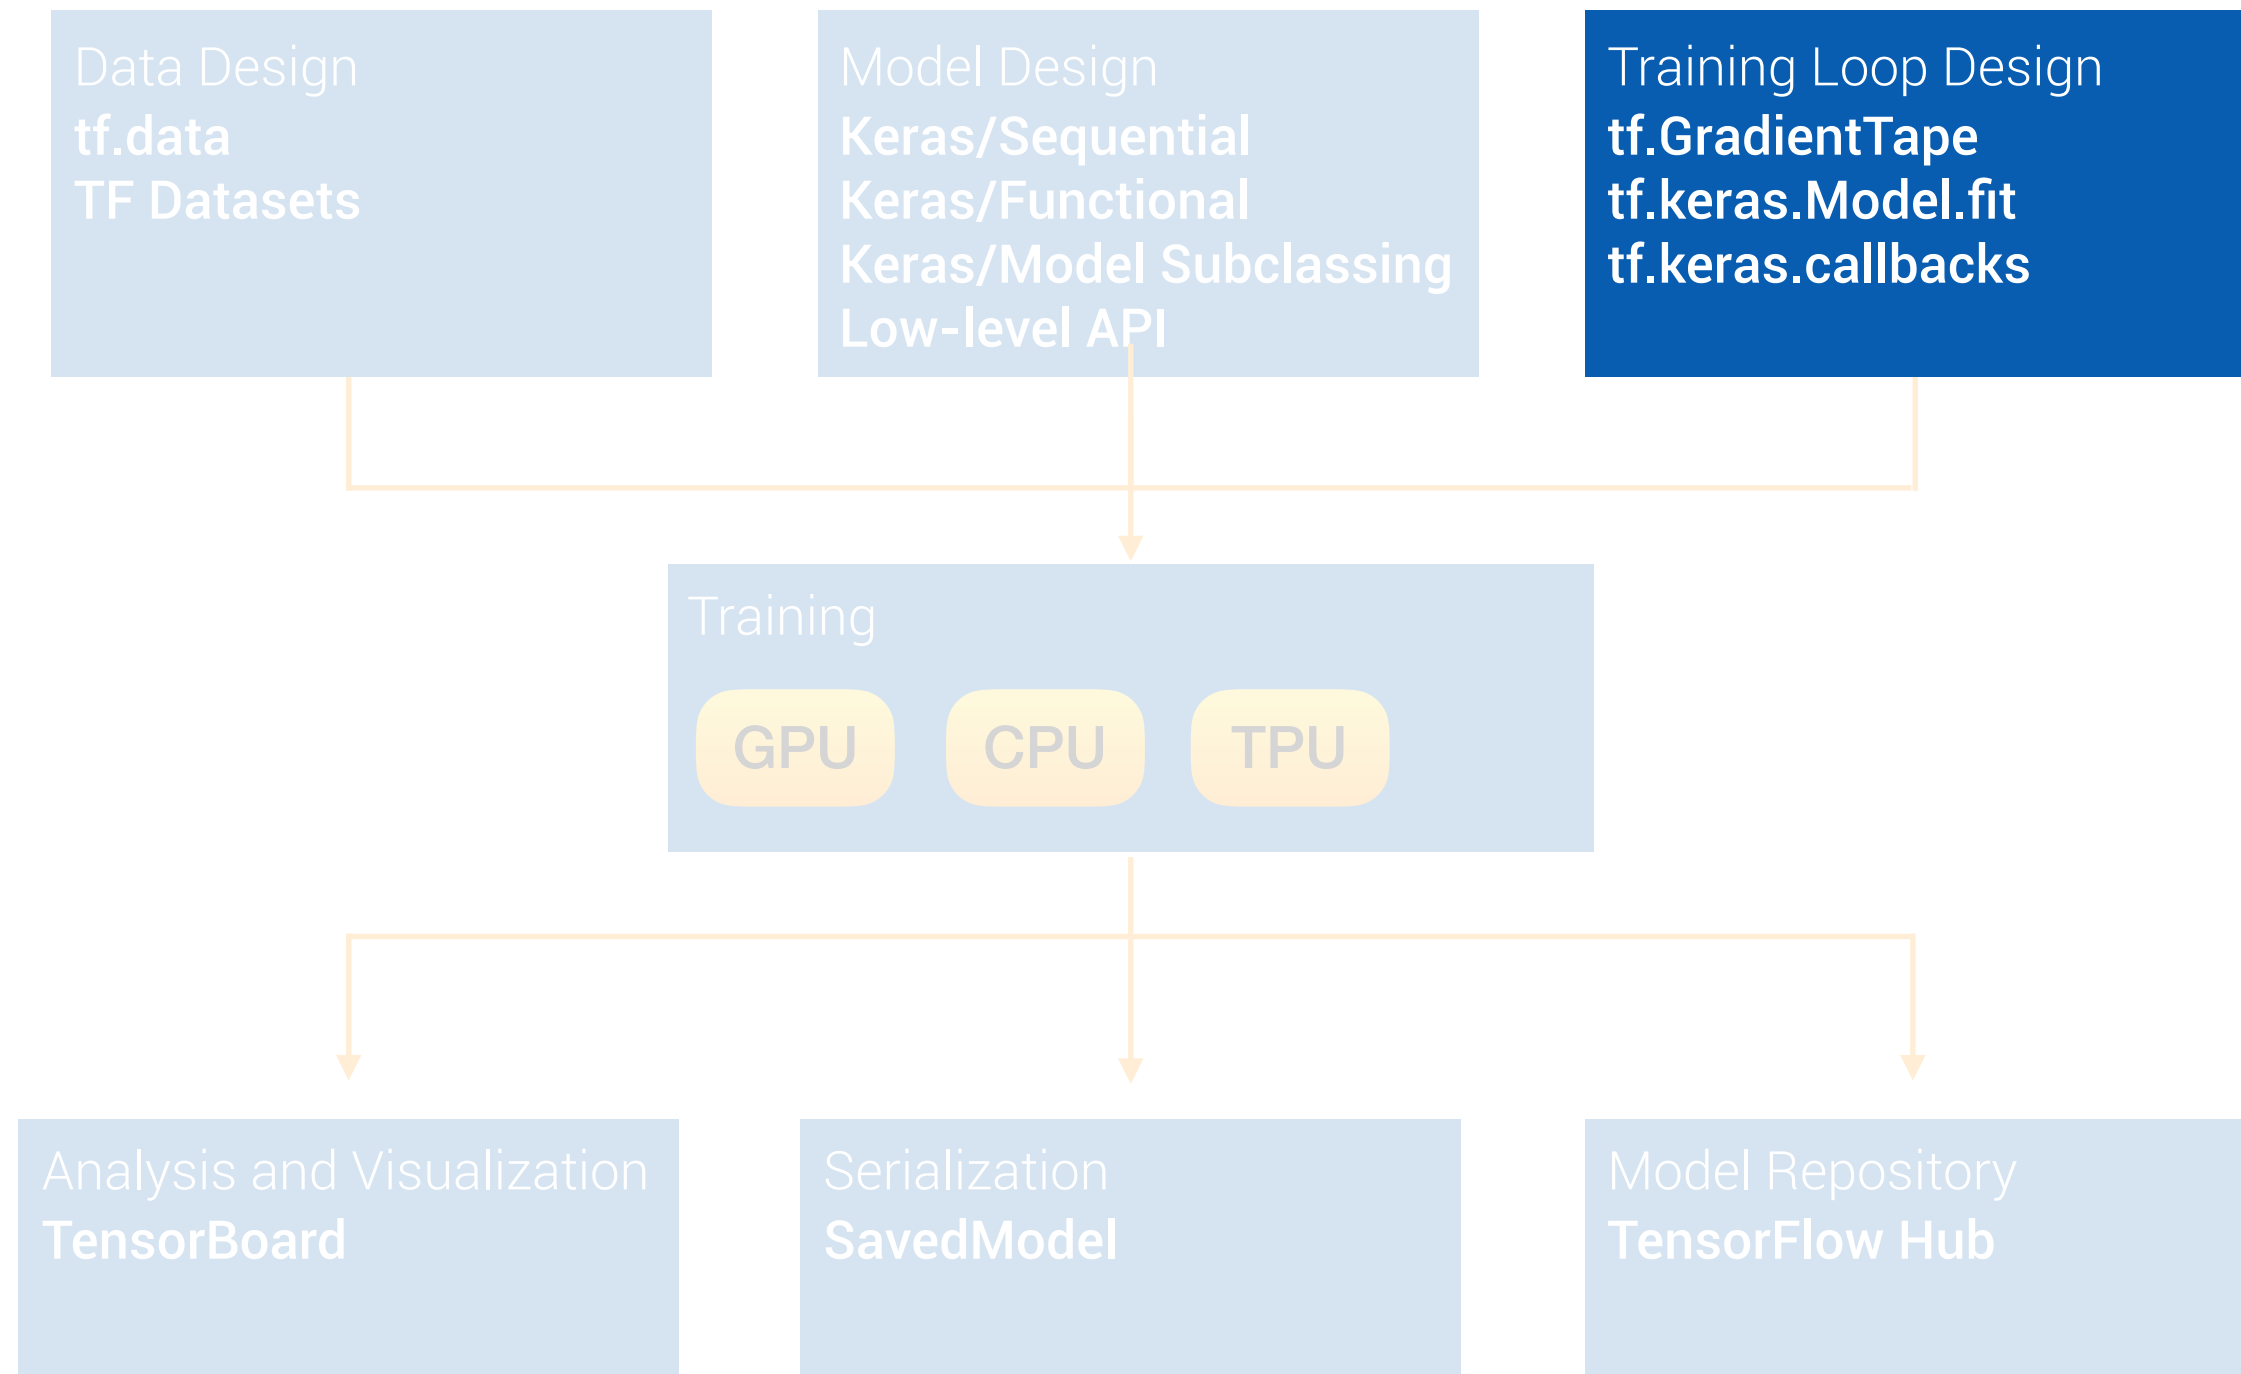
github.com/iust-deep-learning/tensorflow-2-tutorial/tree/master/part_04_model_training_apis



TensorFlow Overview



TensorFlow Overview



Model Training



Model Training

- Keras **built-in** loops

- model.fit()
+ callbacks

- Fast prototyping



Keras built-in training loops

```
model = MyModel()
```

Keras built-in training loops

```
model = MyModel()
```

Keras Sequential,
Functional,
and Model Sub-classing APIs

Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

- **.compile()** is about configuring the training process, such as specifying the optimizer, loss, and metrics.

Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

```
history = model.fit(data,  
                    epochs=10, batch_size=128,  
                    validation_data=val_data,  
                    callbacks=[EarlyStopping(),  
                               TensorBoard(),  
                               ModelCheckpoint()])
```

Keras built-in training loops

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()])
```

```
history = model.fit(data,  
                    epochs=10, batch_size=128,  
                    validation_data=val_data,  
                    callbacks=[EarlyStopping(),  
                               TensorBoard(),  
                               ModelCheckpoint()])
```

```
results = model.evaluate(test_data, batch_size=128)
```

Keras built-in training loops (run in dynamic graph mode)

```
model = MyModel()
```

```
model.compile(optimizer=Adam(),  
              loss=BinaryCrossentropy(),  
              metrics=[AUC(), Precision(), Recall()],  
              run_eagerly=True)
```

```
history = model.fit(data,  
                    epochs=10, batch_size=128,  
                    validation_data=val_data,  
                    callbacks=[EarlyStopping(),  
                               TensorBoard(),  
                               ModelCheckpoint()])
```

```
results = model.evaluate(test_data, batch_size=128)
```

```
predictions = model.predict(x_test[:3])
```

Model Training

- Keras **built-in** loops

- model.fit()
+ callbacks

- Fast prototyping



Model Training

- Keras **built-in** loops

- model.fit()
+ callbacks

- Fast prototyping

- **GradientTape**

- custom training loops

- Complete control



Gradient tapes

Tensorflow "records" all operations executed inside the context of a *tf.GradientTape* onto a "tape".

Tensorflow then uses that tape and the gradients associated with each recorded operation to compute the gradients of a "recorded" computation using reverse mode differentiation.

Gradient Tape

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    y = x * x

dy_dx = g.gradient(y, x) # Will compute to 6.0
```

Gradient Tape

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    y = x * x

dy_dx = g.gradient(y, x) # Will compute to 6.0
```

How to calculate second-order derivatives?

Gradient Tape

```
x = tf.constant(3.0)
with tf.GradientTape() as g:
    g.watch(x)
    with tf.GradientTape() as gg:
        gg.watch(x)
        y = x * x
    dy_dx = gg.gradient(y, x)      # Will compute to 6.0

d2y_dx2 = g.gradient(dy_dx, x)  # Will compute to 2.0
```

Gradient Tape; A real-world usage

Gradient Tape; A real-world usage

```
model = MyModel()
```

Keras Sequential,
Functional,
and Model Sub-classing APIs

Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()  
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```
logits = model(features)
```

```
loss = loss_fn(labels, logits)
```


Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```
logits = model(features)
```

```
loss = loss_fn(labels, logits)
```

Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```
with tf.GradientTape() as tape:  
    logits = model(features)  
    loss = loss_fn(labels, logits)
```

Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```
with tf.GradientTape() as tape:
```

```
    logits = model(features)
```

```
    loss = loss_fn(labels, logits)
```

```
grads = tape.gradient(loss, model.trainable_variables)
```

```
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```
with tf.GradientTape() as tape:
```

```
    logits = model(features)
```

```
    loss = loss_fn(labels, logits)
```

```
grads = tape.gradient(loss, model.trainable_variables)
```

```
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

Gradient Tape; A real-world usage

```
model = MyModel()
```

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
```

```
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```
def train_step(features, labels):
```

```
    with tf.GradientTape() as tape:
```

```
        logits = model(features)
```

```
        loss = loss_fn(labels, logits)
```

```
    grads = tape.gradient(loss, model.trainable_variables)
```

```
    optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

Gradient Tape; A real-world usage

```
model = MyModel()

loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)

@tf.function # optional
def train_step(features, labels):
    with tf.GradientTape() as tape:
        logits = model(features)
        loss = loss_fn(labels, logits)

    grads = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

Gradient Tape; A real-world usage

```
model = MyModel()

loss_fn = tf.keras.losses.SparseCategoricalCrossentropy()
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)

@tf.function # optional
def train_step(features, labels):
    with tf.GradientTape() as tape:
        logits = model(features)
        loss = loss_fn(labels, logits)

        grads = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(grads, model.trainable_variables))

    for features, labels in data:
        loss = train_step(features, labels)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

Gradient Tape; A real-world usage: Measuring the Accuracy

```
model = ...; optimizer = ...; loss_fn = ...
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
model = ...; optimizer = ...; loss_fn = ...
```

```
train_loss = tf.keras.metrics.Mean(name='train_loss')
```

```
train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(  
    name='train_accuracy')
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
model = ...; optimizer = ...; loss_fn = ...
```

```
train_loss = tf.keras.metrics.Mean(name='train_loss')  
train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(  
    name='train_accuracy')
```

```
@tf.function
```

```
def train_step(features, labels):
```

```
    with tf.GradientTape() as tape:
```

```
        logits = model(features)
```

```
        loss = loss_fn(labels, logits)
```

```
    grads = tape.gradient(loss, model.trainable_variables)
```

```
    optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
model = ...; optimizer = ...; loss_fn = ...
```

```
train_loss = tf.keras.metrics.Mean(name='train_loss')  
train_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(  
    name='train_accuracy')
```

```
@tf.function
```

```
def train_step(features, labels):
```

```
    with tf.GradientTape() as tape:
```

```
        logits = model(features)
```

```
        loss = loss_fn(labels, logits)
```

```
    grads = tape.gradient(loss, model.trainable_variables)
```

```
    optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

```
    train_loss(loss)
```

```
    train_accuracy(labels, logits)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
for images, labels in train_ds:  
    train_step(images, labels)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
EPOCHS = 5
```

```
for epoch in range(EPOCHS):
```

```
    for images, labels in train_ds:  
        train_step(images, labels)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
EPOCHS = 5
```

```
for epoch in range(EPOCHS):
```

```
    for images, labels in train_ds:  
        train_step(images, labels)
```

```
    template = 'Epoch {}, Loss: {}, Accuracy: {}'  
    print(template.format(epoch+1,  
                           train_loss.result(),  
                           train_accuracy.result()*100))
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
EPOCHS = 5
for epoch in range(EPOCHS):
    # Reset the metrics at the start of the next epoch
    train_loss.reset_states()
    train_accuracy.reset_states()

    for images, labels in train_ds:
        train_step(images, labels)

    template = 'Epoch {}, Loss: {}, Accuracy: {}'
    print(template.format(epoch+1,
                           train_loss.result(),
                           train_accuracy.result()*100))
```


How to use tf.keras.metrics?

```
m = SomeMetric(...)
for input in ...:
    m.update_state(input)

print('Final result: ', m.result())
```

How to use tf.keras.metrics?

```
m = SomeMetric(...)
for input in ...:
    m.update_state(input)

print('Final result: ', m.result())
```

or

```
m = SomeMetric(...)
for input in ...:
    print('Current result': m(input))

print('Final result: ', m.result().numpy())
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
EPOCHS = 5
```

```
for epoch in range(EPOCHS):
```

```
    # Reset the metrics at the start of the next epoch
```

```
    train_loss.reset_states()
```

```
    train_accuracy.reset_states()
```

```
    for step, (images, labels) in enumerate(train_ds):
```

```
        train_step(images, labels)
```

```
template = 'Epoch {}, Loss: {}, Accuracy: {}'
```

```
print(template.format(epoch+1,
```

```
                train_loss.result(),
```

```
                train_accuracy.result()*100)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
EPOCHS = 5
```

```
for epoch in range(EPOCHS):
```

```
    # Reset the metrics at the start of the next epoch
```

```
    train_loss.reset_states()
```

```
    train_accuracy.reset_states()
```

```
    for step, (images, labels) in enumerate(train_ds):
```

```
        train_step(images, labels)
```

```
        if step % 50 == 0:
```

```
            print(f'Step {step}, Loss{train_loss.result()}')
```

```
template = 'Epoch {}, Loss: {}, Accuracy: {}'
```

```
print(template.format(epoch+1,
```

```
                    train_loss.result(),
```

```
                    train_accuracy.result()*100)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

...

```
for step, (images, labels) in enumerate(train_ds):  
    train_step(images, labels)  
  
    if step % 50 == 0:  
        print(f'Step {step}, Loss{train_loss.result()}')
```

...

Gradient Tape; A real-world usage: Measuring the Accuracy

...

```
for step, (images, labels) in enumerate(train_ds):  
    train_step(images, labels)  
  
    if step % 50 == 0:  
        print(f'Step {step}, Loss{train_loss.result()}')
```

...

Gradient Tape; A real-world usage: Measuring the Accuracy

...

```
for step, (images, labels) in enumerate(train_ds):
    train_step(images, labels)

    if step % 50 == 0:
        print(f'Step {step}, Loss{train_loss.result()}')

    if step % 500 == 0:
        for valid_images, valid_labels in valid_ds:
            test_step(valid_images, valid_labels)

        print(f'Step {step}',
              f'Valid Loss: {valid_loss.result()}',
              f'Valid Accuracy: {valid_accuracy.result()*100}')
```

...

Gradient Tape; A real-world usage: Measuring the Accuracy

...

```
for step, (images, labels) in enumerate(train_ds):
    train_step(images, labels)

    if step % 50 == 0:
        print(f'Step {step}, Loss{train_loss.result()}')

    if step % 500 == 0:
        for valid_images, valid_labels in valid_ds:
            test_step(valid_images, valid_labels)

        print(f'Step {step}',
              f'Valid Loss: {valid_loss.result()}',
              f'Valid Accuracy: {valid_accuracy.result()*100}')
```

...

Gradient Tape; A real-world usage: Measuring the Accuracy

```
@tf.function
def test_step(features, labels):
    logits = model(images)
    loss = loss_object(labels, logits)

    valid_loss(loss)
    valid_accuracy(labels, predictions)
```

Gradient Tape; A real-world usage: Measuring the Accuracy

```
valid_loss = tf.keras.metrics.Mean(name='valid_loss')  
valid_accuracy = tf.keras.metrics.SparseCategoricalAccuracy(  
    name='valid_accuracy')
```

```
@tf.function  
def test_step(features, labels):  
    logits = model(images)  
    loss = loss_object(labels, logits)  
  
    valid_loss(loss)  
    valid_accuracy(labels, predictions)
```

Summary

- ▶ Keras built-in training loops
- ▶ TF Gradient Tape
- ▶ Custom Training Loops
- ▶ Using `tf.keras.metrics`

Thank you!

